

# Computer Vision & Image Processing CSE 473 / 573

Instructor - Kevin R. Keane, PhD

TAs - Radhakrishna Dasari, Yuhao Du, Niyazi Sorkunlu


Lecture 4 – September 6, 2017

# Physical parameters of image formation

- Photometric
  - Type, direction, intensity of light reaching sensor
  - Surfaces' reflectance properties
  - Inference from shading

# Image formation

- What determines the brightness of an image pixel?



# Image formation - photometry


- The brightness of a pixel is a function of the brightness of the surface in the scene that projects to the pixel
- The brightness of the pixel is dependent on
  - How much light is incident on the surface and
  - The fraction of incident light that gets reflected
- We will now explore a few simple models of shading

# Key issues

- Physical
  - what makes a pixel take its brightness values?
- Inference
  - what can we recover from the world using those brightness values?
- Human
  - What can people do?
    - which suggests problems we might be able to solve in computer vision

# The interaction of light and surfaces

- What happens when a light ray hits a point on an object?
  - Some of the light gets **absorbed**
    - converted to other forms of energy (e.g., heat)
  - Some gets **transmitted** through the object
    - possibly bent, through refraction
    - or scattered inside the object (subsurface scattering)
  - Some gets **reflected**
    - possibly in multiple directions at once
  - Really complicated things can happen
    - fluorescence



Source: Steve Seitz

# Model

- Light arrives at a surface
  - from a light source
  - from another surface
- It is reflected into the camera
  - many possible effects
- It arrives at a sensor at the back of the camera
  - a record is made
  - this could be a linear or a non-linear function of the amount of light

From light rays to pixel values

Pixel brightness depends on:

- Camera response
- Surface reflection
- Illumination




# From light rays to pixel values

## Pixel brightness depends on:

- Camera response
- Surface reflection
- Illumination


# Camera response

- Film
  - Record is made by chemical processes in the film
  - These are non-linear; typically,
    - dark patches are lighter than they should be
    - light patches are darker than they should be
    - these ensure that detail is visible



# Camera response

- CCD
  - Linear devices
    - with non-linearities produced by electronics to mimic film
- Calibration
  - Can be hard to find curves of camera response vs light input
  - Instead, use calibration algorithms (no covered in our course)



# From light rays to pixel values

## Pixel brightness depends on:

- Camera response
- Surface reflection
- Illumination


# Surface reflection

- Many effects can occur when light strikes a surface -- could be:
  - absorbed; transmitted; reflected; scattered
    - e.g. some people can see arteries, veins under their skin
      - because light is transmitted through skin, reflected at blood vessel, transmitted out
- For simplicity, we assume that
  - surfaces don't fluoresce
  - surfaces don't emit light (i.e. are cool)
  - all the light leaving a point is due to that arriving at that point

# The important reflection modes



## 1. Specular reflection (mirror like)

- Pure mirror:
  - incoming, outgoing directions and normal are coplanar
  - incoming, outgoing angles to normal are equal
- Most specular surfaces:
  - some light leaves the surface along directions near to the specular direction as well




## 2. Diffuse reflection

- Light leaves in equal amounts in each direction
  - so surface looks equally bright from each viewing direction



# Reflection modes



- The 2 most important reflection modes are
  - Diffuse reflection – incident light is spread evenly over the whole hemisphere of out going directions
  - Specular reflection – reflected light is concentrated in a single direction
    - Specular direction  $S$  is coplanar with the normal  $N$  and source direction  $L$
    - Incident angle = reflection angle =  $\theta$


# Diffuse reflection

- Light leaves the surface evenly in all directions
  - e.g. cotton cloth, carpets, matte paper, matte paints, most “rough” surfaces
- Described by one parameter: **albedo**
  - percentage of light arriving that leaves
  - range 0-1 practical range is smaller
- Light leaving is *albedo x light arriving*
  - Ambiguity: A surface could be dark because
    - It reflects a small percentage of the light arriving
    - There isn't very much light arriving



# Diffuse reflection


- Light is reflected equally in all directions
  - Dull, matte surfaces like chalk or latex paint
  - Microfacets scatter incoming light randomly
  - Effect is that light is reflected equally in all directions
- Brightness of the surface depends on the incidence of illumination



brighter





darker




# How much light arrives?

- Assume source is far away
  - So light travels in parallel rays
  - (Light arriving) is proportional to (number of rays striking surface)
  - Surface A below receives less light than surface B
- Drawing yields
  - (number of rays striking surface) proportional to  $\cos(\theta)$ 
    - where  $\theta$  is angle between normal and direction of travel
- Shadows
  - If point can't see the light source, it is in shadow




# Diffuse reflection: Lambert's law



$$\begin{aligned} B &= \rho (\mathbf{N} \cdot \mathbf{S}) \\ &= \rho \|\mathbf{S}\| \cos \theta \end{aligned}$$


$B$ : radiosity (total power leaving the surface per unit area)  
 $\rho$ : albedo (fraction of incident irradiance reflected by the surface)  
 $N$ : unit normal  
 $S$ : source vector (magnitude proportional to intensity of the source)




# Specular reflection

- Radiation arriving along a source direction leaves along the specular direction (source direction reflected about normal)
- Some fraction is absorbed, some reflected
- On real surfaces, energy usually goes into a lobe of directions
- Phong model: reflected energy falls off with  $\cos^n(\delta\theta)$
- Lambertian + specular model: sum of diffuse and specular term


$$\cos^n(\delta\theta)$$



# Specular reflection



Moving the light source



Changing the exponent

# Diffuse + Specular model

- Most surfaces can be modeled as diffuse + specular
  - surface parameters:
    - diffuse albedo,
    - specular albedo,
- This justifies the following strategy for many analyses
  - Find and remove specularities
    - which are small, and bright
    - More sophisticated specularity finders use color information
  - Ignore the missing points, and treat the rest of the image as diffuse

# Diffuse + Specular model cont'd

Choose the source effects to model:

- Simplest case is the local shading model
  - assume light comes from the source, no interreflections
  - Assume source is at an infinitely distance
- Again we are working with  $N(x)$  and  $S(x)$  and  $p(x)$  is albedo at  $x$ .
- Let  $Vis(S, x)$  be a function that is 1 when  $x$  can see the source, and 0 otherwise
- Then the intensity at  $x$  is given as:

$$I(\mathbf{x}) = \rho(\mathbf{x})(\mathbf{N} \cdot \mathbf{S})Vis(S, \mathbf{x}) + \rho(\mathbf{x})A + M$$

Image Intensity		Diffuse term		Ambient term		Specular mirror-like term
--------------------	--	--------------	--	-----------------	--	---------------------------------

# From light rays to pixel values


## Pixel brightness depends on:

- Camera response
- Surface reflection
- Illumination





# Shadows

- Most shadows aren't dark
  - because shadow points get light from other surfaces, not just light source
- Area sources
  - Large, bright areas
  - e.g. diffuser boxes, the sky
  - Yield smooth, blurry shadows
    - Points that can see the whole source are brighter
    - Points that can see only part of the source are darker (penumbra)
    - Points that can see no part of the source are darkest (umbra)
- Other surfaces behave like area sources
  - Smooth, blurry shadows are common (and sometimes too faint to see)




# How much light arrives?



- Orientation of surface affects how much light the surface gathers
- Patch A is tilted from the source with  $\theta$  close to  $90^\circ$  collects less energy (it cuts fewer rays per unit surface area)
- Patch B is facing the source with  $\theta$  close to  $0^\circ$  so it appears brighter
- The darker surfaces are turned away from the illumination direction

# Area sources



Area sources generate complex shadows with smooth boundaries

- The surface patch sees the source disappearing slowly behind the occluder
- Patch A is darker than B
  - Shadowing effect but no clear boundary (instead is a smooth gradient)
- Chair leg casts a complex shadow with 2 distinct regions
  - **Umbra** – no source is seen at all
  - **Penumbra** – source is partially seen

# Light and shading - Crucial points

- Image brightness is affected by
  - amount of light arriving at surface
  - surface type (diffuse, specular) and amount reflected at surface
  - camera sensitivity
- There are significant ambiguities
  - eg low albedo surface in bright light
    - vs high albedo surface in low light
  - each might reflect about the same amount
- Most surfaces can be modeled as diffuse + specular
  - generally, find and remove specularities
  - treat the rest as diffuse
- Shadows occur when a light source is obstructed by an object.
  - occupies all of the space behind the opaque object with light in front of it.

# Inference from shading

- Radiometric calibration and high dynamic range images
- The shape of specularities
- Inferring lightness and illumination
- Photometric stereo: shape from multiple shaded images


# Photometric stereo (shape from shading)

- Can we reconstruct the shape of an object based on shading cues?



Luca della Robbia,  
*Cantoria*, 1438

# Surface model: Monge patch



A Monge patch is a representation of a piece of surface as a height function. An orthographic camera, that maps  $(x, y, z)$  to  $(x, y)$  in the camera is viewing a Monge patch, i.e. the shape of the surface is represented a function of position in the image

# Surface model continued

- Photometric stereo is a method for recovering a representation of the Monge patch from image data
  - Reason about image intensity for different images of the surface under different illumination conditions
  - Recover height of surface at each image pixel
  - Fix the positions of the camera and surface
    - Illuminate using faraway source (local shading model)
- The intensity value of a pixel at  $(x,y)$  is now

$$B(\mathbf{x}) = \rho(\mathbf{x})\mathbf{N}(\mathbf{x})\mathbf{S}(\mathbf{x})$$

$$I(x, y) = kB(\mathbf{x})$$



# Image model

- **Known:** source vectors  $\mathbf{S}_j$  and pixel values  $I_j(x,y)$ 
  - $j$  is the index of the illumination source
- **Unknown:** surface normal  $\mathbf{N}(x,y)$  and albedo  $\rho(x,y)$
- Assume that the response function of the camera is a linear scaling by a factor of  $k$

- Lambert's law:
$$\begin{aligned} I_j(x, y) &= kB(\mathbf{x}) \\ &= k\rho(x, y)(\mathbf{N}(x, y) \cdot \mathbf{S}_j) \\ &= (\rho(x, y)\mathbf{N}(x, y)) \cdot (k\mathbf{S}_j) \\ &= \mathbf{g}(x, y) \cdot \mathbf{V}_j \end{aligned}$$

## Image model cont'd


- $g(x,y)$  describes the surface
- $V_j$  is a property of the illumination and camera
- We can have a dot product between the vector field  $g(x,y)$  and the vector  $V_j$
- So for  $n$  sources (if  $n$  is sufficiently large), we can stack up the known  $V_j$  into a matrix  $V$
- For each point in the image, we stack up the intensity measurement from each of the  $n$  sources

# Least squares problem


- For each pixel, set up a linear system:
- Obtain least-squares solution for  $\mathbf{g}(x,y)$  (which we defined as  $\mathbf{N}(x,y) \rho(x,y)$ )
- Since  $\mathbf{N}(x,y)$  is the unit normal,  $\rho(x,y)$  is given by the magnitude of  $\mathbf{g}(x,y)$
- Finally,  $\mathbf{N}(x,y) = \mathbf{g}(x,y) / \rho(x,y)$

$$\begin{array}{ccc}
 \left[ \begin{array}{c} I_1 ( x , y ) \\ I_2 ( x , y ) \\ \vdots \\ I_n ( x , y ) \end{array} \right] & = & \left[ \begin{array}{c} \mathbf{V}_1^T \\ \mathbf{V}_2^T \\ \vdots \\ \mathbf{V}_n^T \end{array} \right] \mathbf{g} ( x , y ) \\
 \begin{array}{c} | \\ (n \times 1) \\ \text{known} \end{array} & & \begin{array}{cc} | & | \\ (n \times 3) & (3 \times 1) \\ \text{known} & \text{unknown} \end{array}
 \end{array}$$


# Example



Recovered albedo




Recovered normal field



F&P 2<sup>nd</sup> ed., sec. 2.2.4

# Surface recovered by integration



# Limitations

- Orthographic camera model
- Simplistic reflectance and lighting model
- No shadows
- No inter-reflections
- No missing data
- Integration is tricky

# Computer vision application


Finding the direction of light source



P. Nillius and J.-O. Eklundh, "Automatic estimation of the projected light source direction," CVPR 2001

# Computer vision application

Detecting composite photos:



Fake photo




Real photo

M. K. Johnson and H. Farid, [Exposing Digital Forgeries by Detecting Inconsistencies in Lighting](#), ACM Multimedia and Security Workshop, 2005.



# Computer vision application

Estimating natural illumination from a single outdoor image:



From a single image (left), they estimate the most likely sky appearance (middle) and insert a 3-D object (right). Illumination estimation was done entirely automatically

J-F. Lalonde, A. A. Efros and S G. Narasimhan, [Estimating Natural Illumination from a Single Outdoor Image](#), International Journal on Computer Vision, 98(2):123-145, 2012.

# Slide Credits

- David A. Forsyth - UIUC
- Svetlana Lazebnik – UIUC

# Next class

- Color
- Readings for next lecture:
  - Forsyth and Ponce Chapter 3; Szeliski 2.3.2 (optional)
- Readings for today:
  - Forsyth and Ponce 2.1, 2.2.4; Szeliski 2.2 (optional)

# Questions

