

Computer Vision & Image Processing CSE 473 / 573

Instructor - Kevin R. Keane, PhD

TAs - Radhakrishna Dasari, Yuhao Du, Niyazi Sorkunlu

Lecture 11


September 22, 2017

Local Features

Schedule

- Last class
 - We started local features
- Today
 - More on local features
- Readings for today: Forsyth and Ponce Chapter 5

A hard feature matching problem



NASA Mars Rover images

Overview


- Eigenvector / eigenvalue review
- Corners (Harris Detector)
- Blobs
- Descriptors

Eigenvectors / eigenvalues


- *Eigen*: German word for *characteristic*
- *Eigen-vector*: characteristic direction
- *Eigen-value*: characteristic magnitude
- See the Wall et.al. in our Linear Algebra folder
 - Huge $M \times N$ matrices
 - characteristic column response
 - characteristic row response

Corner Detection: Basic Idea


- small window, shifting in any direction should give a ***large change in intensity***
- Eigenvector: direction
- Eigenvalue: magnitude



“flat” region:
no change in
all directions



“edge”:
no change along
edge direction



“corner”:
big change in all
directions

Harris corner detector summary


- Good corners
 - High contrast
 - Sharp change in edge orientation
- Image features at good corners
 - Large gradients that change direction sharply
 - Will have 2 large eigenvalues
- Compute matrix H by summing over window

$$\mathcal{H} = \sum_{window} \{(\nabla I)(\nabla I)^T\}$$
$$\approx \sum_{window} \left\{ \begin{pmatrix} (\frac{\partial G_{\sigma}}{\partial x} ** I)(\frac{\partial G_{\sigma}}{\partial x} ** I) & (\frac{\partial G_{\sigma}}{\partial x} ** I)(\frac{\partial G_{\sigma}}{\partial y} ** I) \\ (\frac{\partial G_{\sigma}}{\partial x} ** I)(\frac{\partial G_{\sigma}}{\partial y} ** I) & (\frac{\partial G_{\sigma}}{\partial y} ** I)(\frac{\partial G_{\sigma}}{\partial y} ** I) \end{pmatrix} \right\}$$

Overview


- Corners (Harris Detector)
- Blobs
- Descriptors

Blob detection with scale selection




Achieving scale covariance


- Goal: independently detect corresponding regions in scaled versions of the same image
- Need *scale selection* mechanism for finding characteristic region size that is *covariant* with the image transformation



Recall: Edge detection

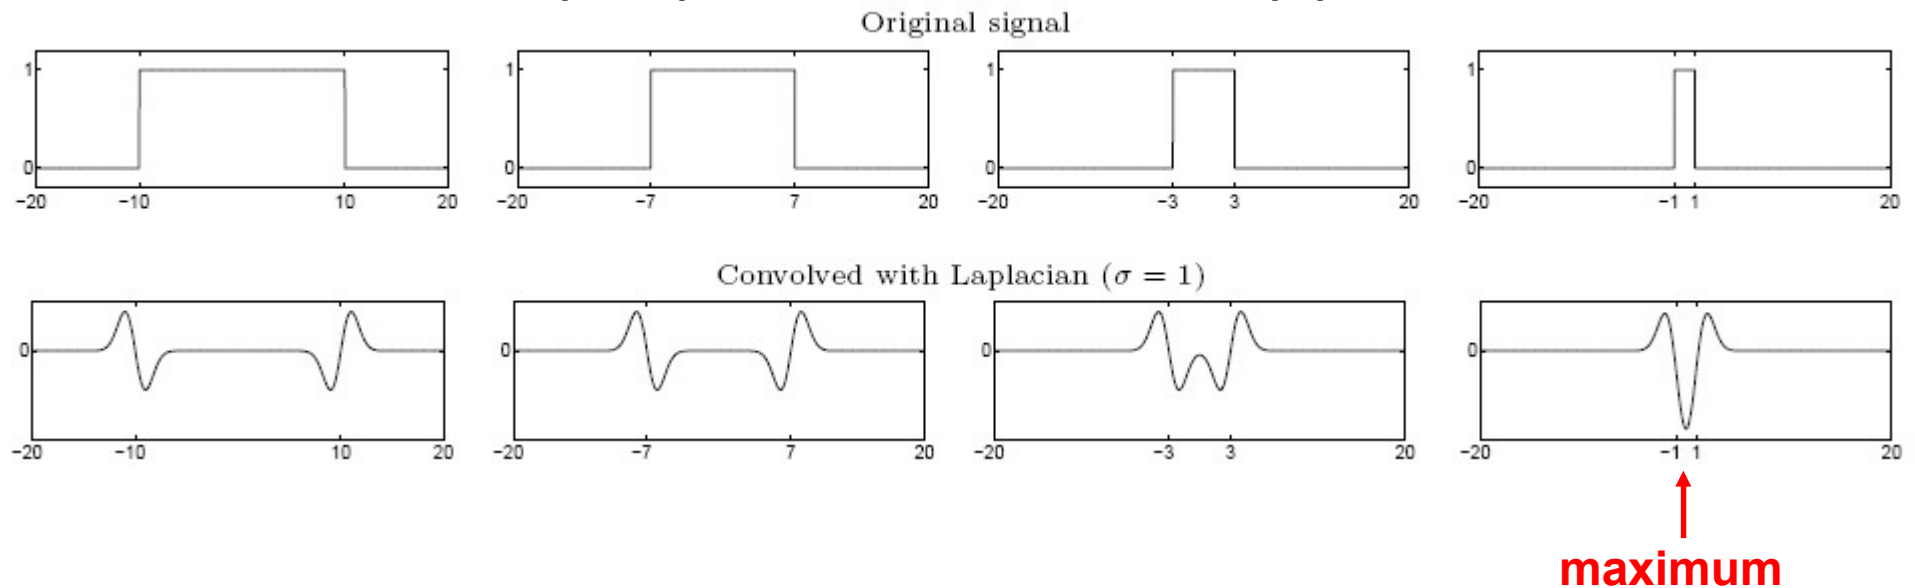


Edge detection, Take 2



From edges to blobs


- Edge = ripple
- Blob = superposition of two ripples



Spatial selection: the magnitude of the Laplacian response will achieve a maximum at the center of the blob, provided the scale of the Laplacian is “matched” to the scale of the blob


Estimating scale - I

- Assume we have detected a corner
- How big is the neighborhood?
- Use Laplacian of Gaussian filter
 - Details on next slide
 - Kernel looks like fuzzy dark blob on pale light foreground
 - Scale (sigma) of Gaussian gives size of dark, light blob
- Strategy
 - Apply Laplacian of Gaussian at different scales at corner
 - response is a function of scale
 - Choose the scale that gives the largest response
 - the scale at which the neighborhood looks “most like” a fuzzy blob
 - This is covariant



Scale selection


- We want to find the characteristic scale of the blob by convolving it with Laplacians at several scales and looking for the maximum response
- However, Laplacian response decays as scale increases:



Why does this happen?

Scale normalization

- The response of a derivative of Gaussian filter to a perfect step edge decreases as σ increases




Scale normalization






- The response of a derivative of Gaussian filter to a perfect step edge decreases as σ increases
- To keep response the same (scale-invariant), must multiply Gaussian derivative by σ
- Laplacian is the second Gaussian derivative, so it must be multiplied by σ^2

Effect of scale normalization






Original signal



Unnormalized Laplacian response





Scale-normalized Laplacian response



↑
maximum

Blob detection in 2D



- Laplacian of Gaussian: Circularly symmetric operator for blob detection in 2D



$$\nabla^2 g = \frac{\partial^2 g}{\partial x^2} + \frac{\partial^2 g}{\partial y^2} = \frac{1}{\pi \sigma^4} \left(1 - \frac{x^2 + y^2}{2\sigma^2} \right) e^{-\frac{x^2 + y^2}{2\sigma^2}}$$

Blob detection in 2D


- Laplacian of Gaussian: Circularly symmetric operator for blob detection in 2D




Scale-normalized:
$$\nabla_{\text{norm}}^2 g = \sigma^2 \left(\frac{\partial^2 g}{\partial x^2} + \frac{\partial^2 g}{\partial y^2} \right)$$

Scale selection


- At what scale does the Laplacian achieve a maximum response to a binary circle of radius r ?



image



Laplacian




Scale selection


- At what scale does the Laplacian achieve a maximum response to a binary circle of radius r ?
- To get maximum response, the zeros of the Laplacian have to be aligned with the circle

- Zeros of Laplacian is given by (up to scale): $\left(1 - \frac{x^2 + y^2}{2\sigma^2}\right) = 0$

- Therefore, the maximum response occurs at $\sigma = r / \sqrt{2}$.




image



Characteristic scale


- We define the characteristic scale of a blob as the scale that produces peak of Laplacian response in the blob center



T. Lindeberg (1998). ["Feature detection with automatic scale selection."](#)
International Journal of Computer Vision **30** (2): pp 77--116.

Scale-space blob detector

1. Convolve image with scale-normalized Laplacian at several scales
2. Find maxima of squared Laplacian response in scale-space



Scale-space blob detector: Example




Scale-space blob detector: Example



sigma = 11.9912

Scale-space blob detector: Example



Efficient implementation


- Approximating the Laplacian with a difference of Gaussians:

$$L = \sigma^2 \left(G_{xx}(x, y, \sigma) + G_{yy}(x, y, \sigma) \right)$$

(Laplacian)

$$DoG = G(x, y, k\sigma) - G(x, y, \sigma)$$

(Difference of Gaussians)



Invariance and covariance properties

- Laplacian (blob) response is *invariant* w.r.t. rotation and scaling
- Blob location is *covariant* w.r.t. rotation and scaling

Estimating scale - summary

- Assume we have detected a corner
- How big is the neighborhood?
- Use Laplacian of Gaussian filter
 - Details on next slide
 - Kernel looks like fuzzy dark blob on pale light foreground
 - Scale (sigma) of Gaussian gives size of dark, light blob
- Strategy
 - Apply Laplacian of Gaussian at different scales at corner
 - response is a function of scale
 - Choose the scale that gives the largest response
 - the scale at which the neighborhood looks “most like” a fuzzy blob
 - This is covariant



Estimating scale - summary

- Laplacian of a function

$$\nabla^2 f = \frac{\partial^2 f}{\partial x^2} + \frac{\partial^2 f}{\partial y^2}$$

- Gaussian $G_\sigma(x, y) = \frac{1}{2\pi\sigma^2} e^{\left(\frac{-x^2 - y^2}{2\sigma^2}\right)}$

- So Laplacian of Gaussian

$$\nabla^2 G_\sigma(x, y) = \left(\frac{x^2 + y^2 - 2\sigma^2}{\sigma^4} \right) G_\sigma(x, y)$$

- Convolve with image

$$\nabla_\sigma^2 \mathcal{I}(x, y) = (\nabla^2 G_\sigma(x, y)) * * \mathcal{I}(x, y)$$

Overview


- Corners (Harris Detector)
- Blobs
- Descriptors (SIFT)

Scale Invariant Feature Transform

David Lowe IJCV 2004


Basic idea:

- Take 16x16 square window around detected feature
- Compute edge orientation (angle of the gradient - 90°) for each pixel
- Throw out weak edges (threshold gradient magnitude)
- Create histogram of surviving edge orientations



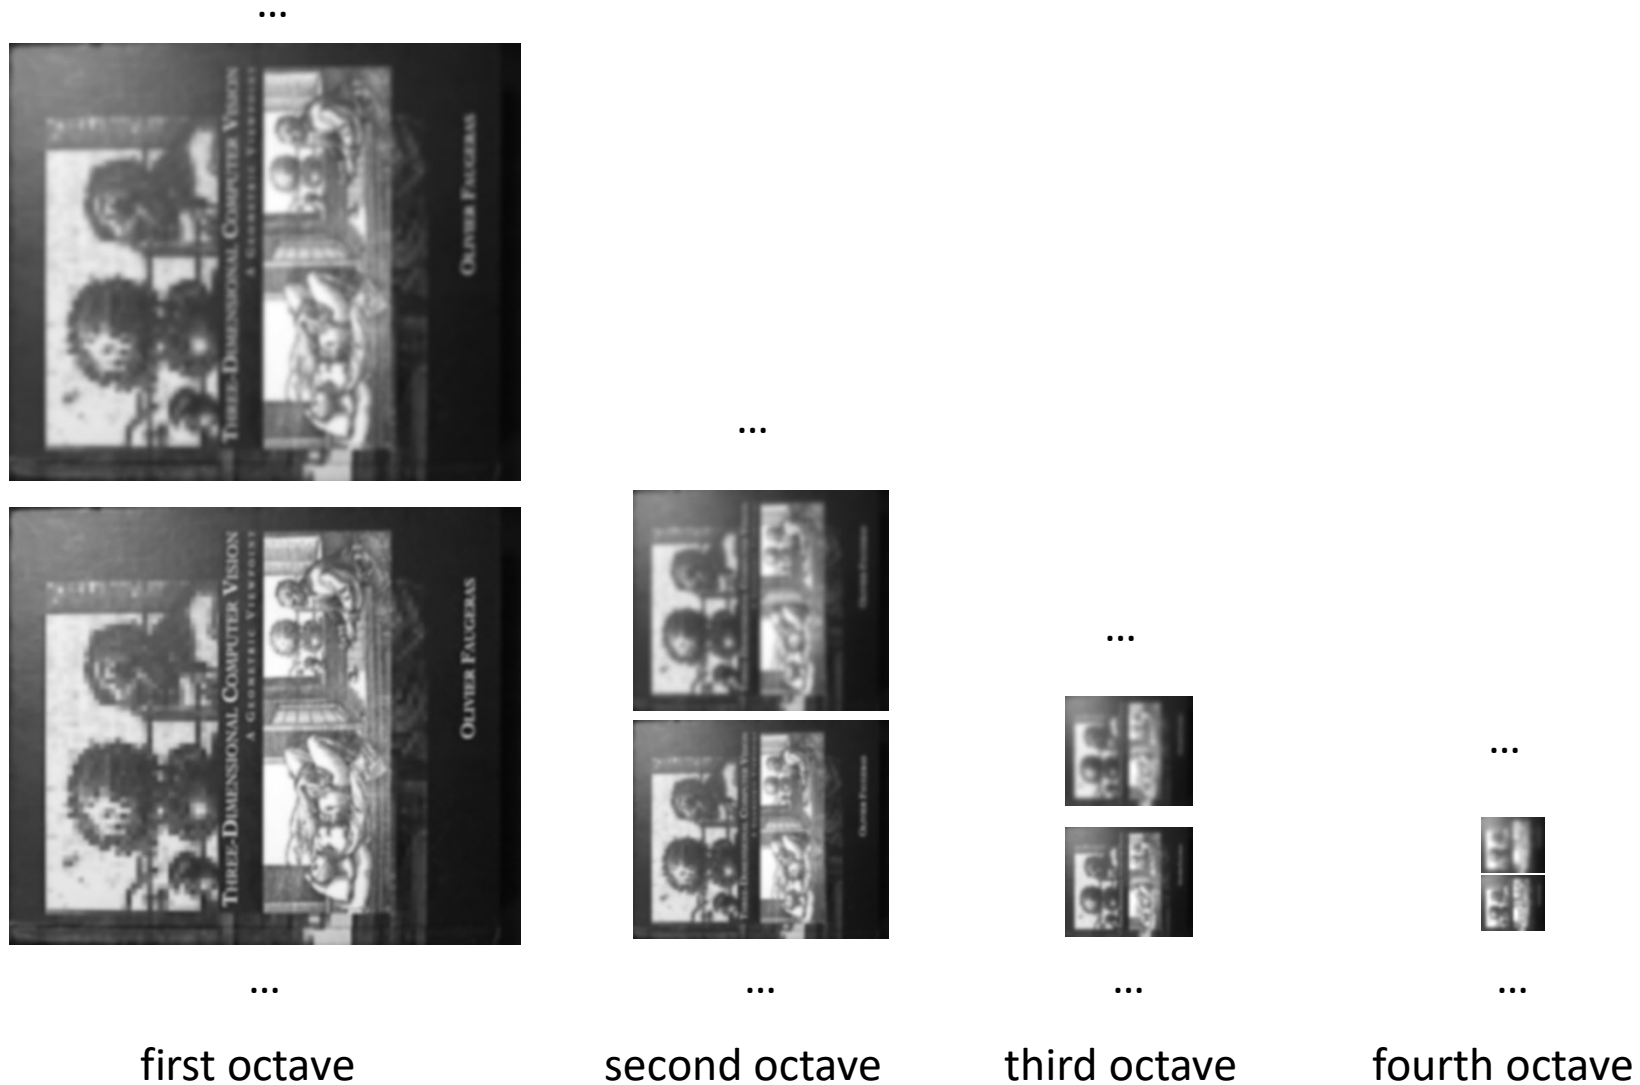
Adapted from slide by David Lowe

Efficient implementation

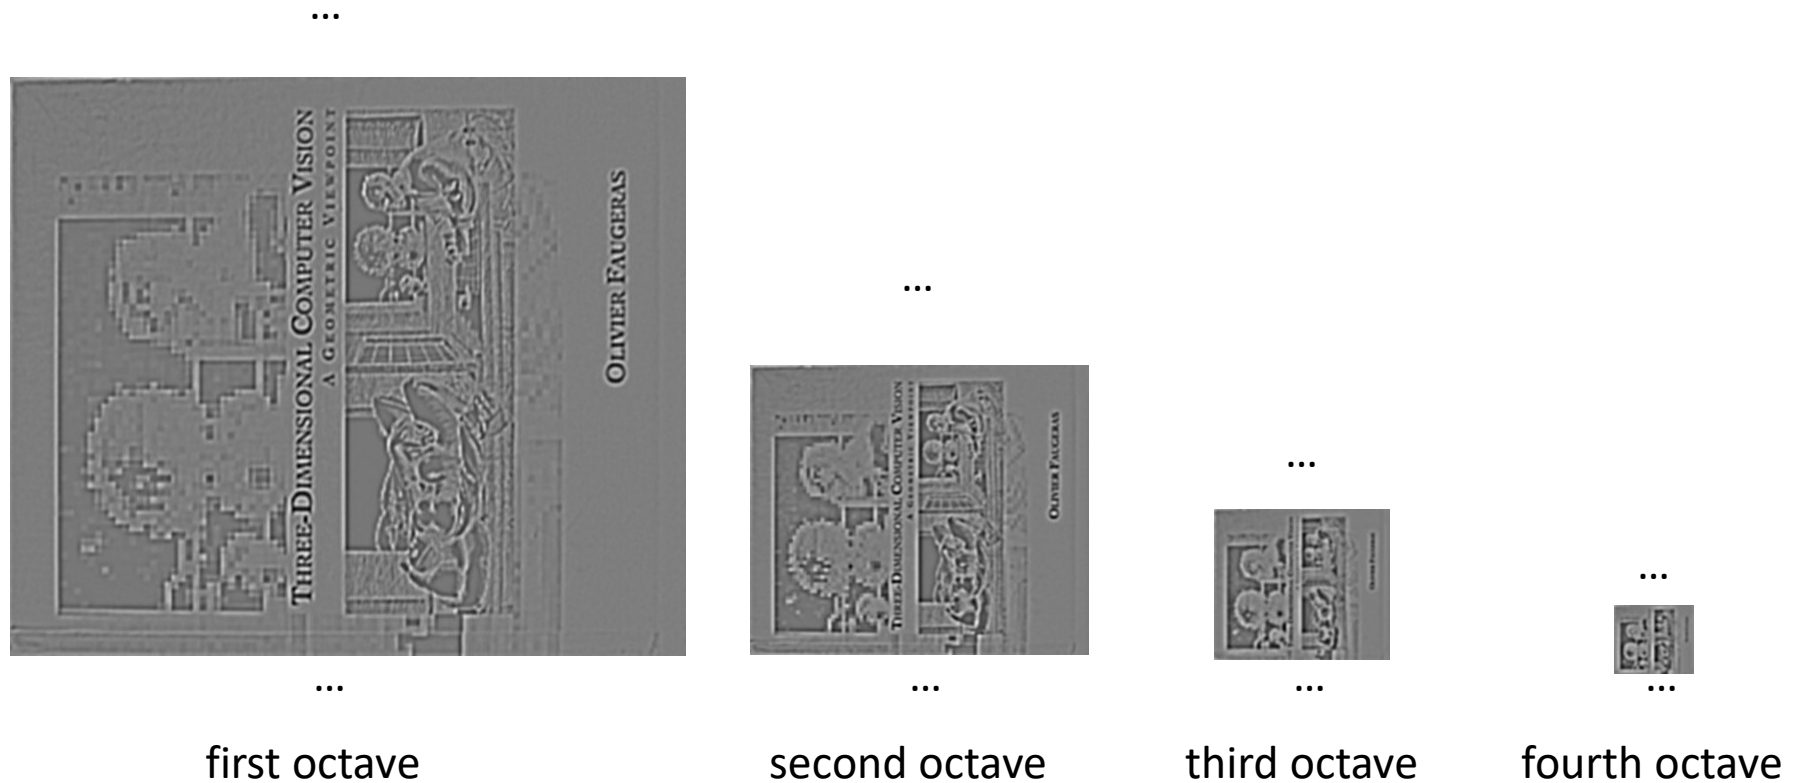


David G. Lowe. ["Distinctive image features from scale-invariant keypoints."](#) *IJCV* 60 (2), pp. 91-110, 2004.

Scale space images




Difference-of-Gaussian images



Finding extrema

- Sample point is selected only if it is a minimum or a maximum of these points



Extrema in this image

Localization

- 3D quadratic function is fit to the local sample points

- Start with Taylor expansion with sample point as the origin

$$D(\mathbf{X}) = D + \frac{\partial D^T}{\partial \mathbf{X}} \mathbf{X} + \frac{1}{2} \mathbf{X}^T \frac{\partial^2 D}{\partial \mathbf{X}^2} \mathbf{X}$$

– where $\mathbf{X} = (x, y, \sigma)^T$

- Take the derivative with respect to \mathbf{X} , and set it to 0, giving

$$0 = \frac{\partial D}{\partial \mathbf{X}} + \frac{\partial^2 D}{\partial \mathbf{X}^2} \hat{\mathbf{X}}$$

- $\hat{\mathbf{X}} = -\frac{\partial^2 D^{-1}}{\partial \mathbf{X}^2} \frac{\partial D}{\partial \mathbf{X}}$ is the location of the keypoint
- This is a 3x3 linear system

Localization

$$\begin{bmatrix} \frac{\partial^2 D}{\partial \sigma^2} & \frac{\partial^2 D}{\partial \sigma y} & \frac{\partial^2 D}{\partial \sigma x} \\ \frac{\partial^2 D}{\partial \sigma y} & \frac{\partial^2 D}{\partial y^2} & \frac{\partial^2 D}{\partial y x} \\ \frac{\partial^2 D}{\partial \sigma x} & \frac{\partial^2 D}{\partial y x} & \frac{\partial^2 D}{\partial x^2} \end{bmatrix} \begin{bmatrix} \sigma \\ y \\ x \end{bmatrix} = - \begin{bmatrix} \frac{\partial D}{\partial \sigma} \\ \frac{\partial D}{\partial y} \\ \frac{\partial D}{\partial x} \end{bmatrix}$$

- Derivatives approximated by finite differences,

– example: $\frac{\partial D}{\partial \sigma} = \frac{D_{k+1}^{i,j} - D_{k-1}^{i,j}}{2}$

$$\frac{\partial^2 D}{\partial \sigma^2} = \frac{D_{k-1}^{i,j} - 2D_k^{i,j} + D_{k+1}^{i,j}}{1}$$

$$\frac{\partial^2 D}{\partial \sigma y} = \frac{(D_{k+1}^{i+1,j} - D_{k-1}^{i+1,j}) - (D_{k+1}^{i-1,j} - D_{k-1}^{i-1,j})}{4}$$

- If X is > 0.5 in any dimension, process repeated

Filtering

- Contrast (use prev. equation): $D(\hat{X}) = D + \frac{1}{2} \frac{\partial D^T}{\partial X} \hat{X}$
 - If $|D(X)| < 0.03$, throw it out
- Edgeiness:
 - Use ratio of principal curvatures to throw out poorly defined peaks
 - Curvatures come from Hessian: $H = \begin{bmatrix} D_{xx} & D_{xy} \\ D_{xy} & D_{yy} \end{bmatrix}$
 - Ratio of $Trace(H)^2$ and $Determinant(H)$
$$Tr(H) = D_{xx} + D_{yy}$$
$$Det(H) = D_{xx}D_{yy} - (D_{xy})^2$$
 - If ratio $> (r+1)^2/(r)$, throw it out (SIFT uses $r=10$)

Orientation assignment


- Descriptor computed relative to keypoint's orientation achieves rotation invariance
- Precomputed along with mag. for all levels (useful in descriptor computation)

$$m(x, y) = \sqrt{(L(x+1, y) - L(x-1, y))^2 + (L(x, y+1) - L(x, y-1))^2}$$

$$\theta(x, y) = a \tan 2((L(x, y+1) - L(x, y-1)) / (L(x+1, y) - L(x-1, y)))$$


- Multiple orientations assigned to keypoints from an orientation histogram
 - Significantly improve stability of matching

Descriptor



- Descriptor has 3 dimensions (x, y, ϑ)
- Orientation histogram of gradient magnitudes
- Position and orientation of each gradient sample rotated relative to keypoint orientation

Descriptor



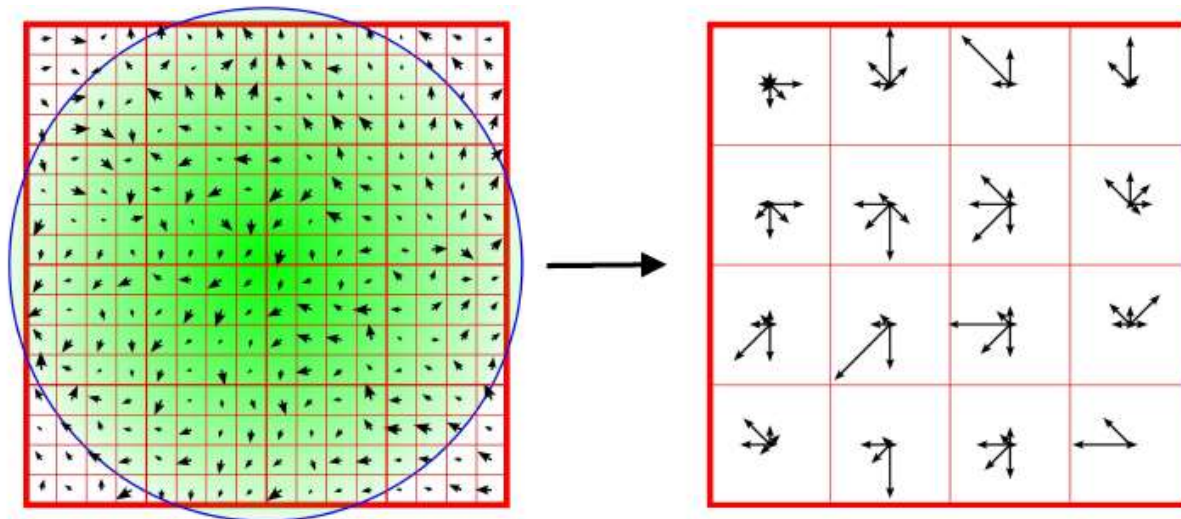
- Weight magnitude of each sample point by Gaussian weighting function
- Distribute each sample to adjacent bins by trilinear interpolation (avoids boundary effects)

Descriptor

- Best results achieved with $4 \times 4 \times 8 = 128$ descriptor size
- Normalize to unit length
 - Reduces effect of illumination change
- Cap each element to 0.2, normalize again
 - Reduces non-linear illumination changes
 - 0.2 determined experimentally

Orientation Histogram

- 4x4 spatial bins (16 bins total)
- Gaussian center-weighting
- 8-bin orientation histogram per bin
- $8 \times 16 = 128$ dimensions total
- Normalized to unit norm



SIFT features



- Very strong record of effectiveness in matching applications
 - use orientations to suppress intensity change effects
 - use histograms so neighborhood need not be exactly localized
 - weight large gradients higher than small gradients
 - Weighting processes are different
 - SIFT features behave very well using nearest neighbors matching
 - i.e. the nearest neighbor to a query patch is usually a matching patch

Scale Invariant Detectors

- Experimental evaluation of detectors w.r.t. scale change

Repeatability rate:

$$\frac{\# \text{ correspondences}}{\# \text{ possible correspondences}}$$



Neighborhoods and SIFT – Key Points

- Algorithms to find neighborhoods
 - Represented by location, scale and orientation
 - Neighborhood is covariant
 - If image is translated, scaled, rotated
 - Neighborhood is translated, scaled, rotated
 - Important property for matching
 - Affine covariant constructions are available
- Once found, describe with SIFT features
 - A representation of local orientation histograms, comparable to HOG
 - Normalized differently

SIFT invariances

- Spatial binning gives tolerance to small shifts in location and scale
- Explicit orientation normalization
- Photometric normalization by making all vectors unit norm
- Orientation histogram gives robustness to small local deformations


PCA-SIFT

- Different descriptor (same keypoints)
- Apply PCA to the gradient patch
- Descriptor size is 20 (instead of 128)
- More robust, faster

Summary of SIFT

Extraordinarily robust matching technique

- Can handle changes in viewpoint
 - Up to about 60 degree out of plane rotation
- Can handle significant changes in illumination
 - Sometimes even day vs. night (below)
- Fast and efficient—can run in real time
- Lots of code available
 - http://people.csail.mit.edu/albert/ladypack/wiki/index.php/Known_implementations_of_SIFT



Slide Credits

- David A. Forsyth - UIUC
- Svetlana Lazebnik – UIUC
- Rob Fergus – NYU

Next class

- Texture
- Readings for next lecture:
 - Forsyth and Ponce 6.1 – 6.4,
 - (optional) Szelinski 10.5
- Readings for today:
 - Forsyth and Ponce 5;
 - (optional) Szelinski 3.1-3.3

Questions

