

Computer Vision & Image Processing CSE 473 / 573

Instructor - Kevin R. Keane, PhD

TAs - Radhakrishna Dasari, Yuhao Du, Niyazi Sorkunlu

Lecture 19

October 13, 2017

Fitting / RANSAC

Slide credits

- Silvio Savarese
- Richard Szeliski

Computer vision

Fitting methods

- Problem formulation
- Least square methods
- RANSAC
- Hough transforms
- Multi-model fitting
- Fitting helps matching!

Fitting

Goal: Choose a parametric model to fit a certain quantity from data

- Lines
- Curves
- Homographic transformation
- Fundamental matrix
- Shape model

Fitting



Goal: Choose a parametric model to fit a certain quantity from data

- Lines
- Curves
- *Homographic transformation*
- Fundamental matrix
- Shape model

Example: Computing vanishing points





Example: Estimating an homographic transformation




H


Example: Estimating F



Example: fitting an 2D shape template



Example: fitting a 3D object model



Fitting

Goal: Choose a parametric model to fit a certain quantity from data


Critical issues:

- noisy data
- outliers
- missing data



Critical issues: noisy data



Critical issues: noisy data (intra-class variability)




Critical issues: outliers



H

Critical issues: missing data (occlusions)



Fitting

Goal: Choose a parametric model to fit a certain quantity from data

Techniques:


- Least square methods
- RANSAC
- Hough transform
- EM (Expectation Maximization)

Least squares methods

- fitting a line -

- **Data:** $(x_1, y_1), \dots, (x_n, y_n)$
- **Line equation:** $y_i = m x_i + b$
- **Find (m, b) to minimize**

$$E = \sum_{i=1}^n (y_i - m x_i - b)^2$$



Least squares methods

- fitting a line -

$$E = \sum_{i=1}^n (y_i - mx_i - b)^2$$

$$E = \sum_{i=1}^n \left(y_i - \begin{bmatrix} x_i & 1 \end{bmatrix} \begin{bmatrix} m \\ b \end{bmatrix} \right)^2 = \left\| \begin{bmatrix} y_1 \\ \vdots \\ y_n \end{bmatrix} - \begin{bmatrix} x_1 & 1 \\ \vdots & \vdots \\ x_n & 1 \end{bmatrix} \begin{bmatrix} m \\ b \end{bmatrix} \right\|^2 = \|Y - XB\|^2$$

$$= (Y - XB)^T (Y - XB) = Y^T Y - 2(XB)^T Y + (XB)^T (XB)$$

$$\frac{dE}{dB} = 2X^T XB - 2X^T Y = 0$$

$$X^T XB = X^T Y$$

Normal equation

$$B = (X^T X)^{-1} X^T Y$$

Least squares methods

- fitting a line -

$$Ax = b$$

- More equations than unknowns
- Look for solution which minimizes $\|Ax-b\| = (Ax-b)^T(Ax-b)$
- Solve $\frac{\partial (Ax-b)^T (Ax-b)}{\partial x_i} = 0$

- LS solution

$$x = (A^T A)^{-1} A^T b$$

Least squares methods

- fitting a line -

Solving $x = (A^t A)^{-1} A^t b$

$$A^+ = (A^t A)^{-1} A^t \quad = \text{pseudo-inverse of } A$$

$$A = U \Sigma V^t \quad = \text{SVD decomposition of } A$$

$$A^{-1} = V \Sigma^{-1} U^t$$

$$A^+ = V \Sigma^+ U^t$$


with Σ^+ equal to Σ^{-1} for all nonzero singular values and zero otherwise

Least squares methods

- fitting a line -

$$E = \sum_{i=1}^n (y_i - mx_i - b)^2$$

$$B = (X^T X)^{-1} X^T Y \quad B = \begin{bmatrix} m \\ b \end{bmatrix}$$



Limitations


- Not rotation-invariant
- Fails completely for vertical lines

Least squares methods

- fitting a line -

- Distance between point (x_n, y_n) and line $ax+by=d$
- Find (a, b, d) to minimize the sum of squared perpendicular distances

$$E = \sum_{i=1}^n (a x_i + b y_i - d)^2$$



Least squares methods

- fitting a line -

$$A h = 0$$


Minimize $\| A h \|$ subject to $\| h \| = 1$

$$A = U D V^T$$

h = last column of V

Least squares methods

- fitting an homography -



$$\begin{bmatrix} wx' \\ wy' \\ w \end{bmatrix} = \begin{bmatrix} h_{11} & h_{12} & h_{13} \\ h_{21} & h_{22} & h_{23} \\ h_{31} & h_{32} & h_{33} \end{bmatrix} \begin{bmatrix} x \\ y \\ 1 \end{bmatrix}$$

Least squares methods

- fitting an homography -

$$h_{11}x + h_{12}y + h_{13} - h_{31}xx' - h_{32}yx' - x' = 0$$


$$h_{21}x + h_{22}y + h_{23} - h_{31}xy' - h_{32}yy' - y' = 0$$

From $n \geq 4$ corresponding points:


$$A h = 0$$

$$\begin{pmatrix} x_1 & y_1 & 1 & 0 & 0 & 0 & -x_1x'_1 & -y_1x'_1 & -x'_1 \\ 0 & 0 & 0 & x_1 & y_1 & 1 & -x_1y'_1 & -y_1y'_1 & -y'_1 \\ x_2 & y_2 & 1 & 0 & 0 & 0 & -x_2x'_2 & -y_2x'_2 & -x'_2 \\ 0 & 0 & 0 & x_2 & y_2 & 1 & -x_2y'_2 & -y_2y'_2 & -y'_2 \\ \vdots & \vdots & \vdots & \vdots & \vdots & \vdots & \vdots & \vdots & \vdots \\ x_n & y_n & 1 & 0 & 0 & 0 & -x_nx'_n & -y_nx'_n & -x'_n \\ 0 & 0 & 0 & x_n & y_n & 1 & -x_ny'_n & -y_ny'_n & -y'_n \end{pmatrix} \begin{bmatrix} h_{1,1} \\ h_{1,2} \\ \vdots \\ h_{3,3} \end{bmatrix} = 0$$

Least squares: Robustness to noise




Least squares: Robustness to noise



Problem: squared error heavily penalizes outliers

Critical issues: outliers




CONCLUSION: Least square is not robust w.r.t. outliers

Least squares: Robust estimators

- General approach:

$$\text{minimize } \sum_i \rho(u_i(x_i, \theta); \sigma) \quad u = \sum_{i=1}^n (y_i - mx_i - b)^2$$


- $u_i(x_i, \theta)$ – residual of i^{th} point w.r.t. model parameters θ
 ρ – robust function with scale parameter σ



The robust function ρ


- favors a configuration with small residuals
- Penalizes large residuals

Least squares: Robust estimators




The effect of the outlier is eliminated

Least squares: Robust estimators



Least squares: Robust estimators



- **CONCLUSION:** Robust estimator useful if prior info about the distribution of points is known
- Robust fitting is a nonlinear optimization problem (iterative solution)
- Least squares solution provides good initial condition

Fitting

Goal: Choose a parametric model to fit a certain quantity from data

Techniques:

- Least square methods
- RANSAC
- Hough transform

Basic philosophy


(voting scheme)

- Data elements are used to vote for one (or multiple) models
- Robust to outliers and missing data
- **Assumption1:** Noise features will not vote consistently for any single model ("few" outliers)
- **Assumption2:** there are enough features to agree on a good model ("few" missing data)

RANSAC

(**RAN**dom **SA**mples **C**onsensus) :
Learning technique to estimate
parameters of a model by random
sampling of observed data

Fischler & Bolles in '81.



$$\pi : I \rightarrow \{P, O\}$$

such that:


$$f(P, \beta) < \delta$$

$$\min_{\pi} |O|$$

Model parameters

$$f(P, \beta) = \left\| \beta - (P^T P)^{-1} P^T \right\|$$

RANSAC




Sample set = set of points in 2D

Algorithm:

1. Select random sample of minimum required size to fit model
 2. Compute a putative model from sample set
 3. Compute the set of inliers to this model from whole data set
- Repeat 1-3 until model with the most inliers over all samples is found

RANSAC




Sample set = set of points in 2D

Algorithm:

1. Select random sample of minimum required size to fit model $[?] = [2]$
 2. Compute a putative model from sample set
 3. Compute the set of inliers to this model from whole data set
- Repeat 1-3 until model with the most inliers over all samples is found

RANSAC




Sample set = set of points in 2D

Algorithm:

1. Select random sample of minimum required size to fit model $[?] = [2]$
 2. Compute a putative model from sample set
 3. Compute the set of inliers to this model from whole data set
- Repeat 1-3 until model with the most inliers over all samples is found

RANSAC



Sample set = set of points in 2D

$$|O| = 14$$

Algorithm:


1. Select random sample of minimum required size to fit model $[?] = [2]$
2. Compute a putative model from sample set
3. Compute the set of inliers to this model from whole data set

Repeat 1-3 until model with the most inliers over all samples is found

RANSAC

(**RAN**dom **SA**mples **C**onsensus) :

Fischler & Bolles in '81.



Algorithm:

1. Select random sample of minimum required size to fit model [?]
 2. Compute a putative model from sample set
 3. Compute the set of inliers to this model from whole data set
- Repeat 1-3 until model with the most inliers over all samples is found

How many samples?


- Number of samples N
 - Choose N so that, with probability p , at least one random sample is free from outliers (e.g. $p=0.99$) (outlier ratio: e)
- Initial number of points s
 - Typically minimum number needed to fit the model
- Distance threshold δ
 - Choose δ so probability for inlier is p (e.g. 0.95)
 - Zero-mean Gaussian noise with std. dev. σ : $t^2=3.84\sigma^2$

$$N = \log(1-p) / \log(1-(1-e)^s)$$

s	proportion of outliers e						
	5%	10%	20%	25%	30%	40%	50%
2	2	3	5	6	7	11	17
3	3	4	7	9	11	19	35
4	3	5	9	13	17	34	72
5	4	6	12	17	26	57	146
6	4	7	16	24	37	97	293
7	4	8	20	33	54	163	588
8	5	9	26	44	78	272	1177

Estimating H by RANSAC

- $H \rightarrow 8$ DOF
- Need 4 correspondences




Sample set = set of matches between 2 images

Algorithm:

1. Select a random sample of minimum required size [?]
 2. Compute a putative model from these
 3. Compute the set of inliers to this model from whole sample space
- Repeat 1-3 until model with the most inliers over all samples is found

Estimating F by RANSAC

- $F \rightarrow 7$ DOF
- Need 7 (8) correspondences




Sample set = set of matches between 2 images


Algorithm:

1. Select a random sample of minimum required size [?]
 2. Compute a putative model from these
 3. Compute the set of inliers to this model from whole sample space
- Repeat 1-3 until model with the most inliers over all samples is found

Matching features



Random Sample Consensus




Fitting helps matching!



Suppose we need to align these 2 images




Fitting helps matching!



Feature are matched (for instance, based on correlation)

Fitting helps matching!



Matches bases on appearance only

Red: good matches

Green: bad matches

Idea:


- Fitting an homography H (by RANSAC) mapping features from images 1 to 2
- Bad matches will be labeled as outliers (hence rejected)!

Fitting helps matching!



Fitting helps matching!





RANSAC - conclusions

Good:

- Simple and easily implementable
- Successful in different contexts

Bad:

- Many parameters to tune
- Trade-off accuracy-vs-time
- Cannot be use if ratio inliers/outliers is too small

Fitting

Goal: Choose a parametric model to fit a certain quantity from data


Techniques:

- Least square methods
- RANSAC
- Hough transform

Hough transform





P.V.C. Hough, *Machine Analysis of Bubble Chamber Pictures*, Proc. Int. Conf. High Energy Accelerators and Instrumentation, 1959

Given a set of , find the curve or line that explains the data points best



$$y = m x + n$$

Hough transform




Hough transform

P.V.C. Hough, *Machine Analysis of Bubble Chamber Pictures*, Proc. Int. Conf. High Energy Accelerators and Instrumentation, 1959


Issue : parameter space $[m,n]$ is unbounded...

- Use a polar representation for the parameter space




$$x \cos \theta + y \sin \theta = \rho$$

Hough transform - experiments




features





votes

Hough transform - experiments



Issue: Grid size needs to be adjusted...

Hough transform - experiments



features

votes

Issue: spurious peaks due to uniform noise

Hough transform - conclusions


Good:

- All points are processed independently, so can cope with occlusion/outliers
- Some robustness to noise: noise points unlikely to contribute consistently to any single bin

Bad:

- Spurious peaks due to uniform noise
- Trade-off noise-grid size (hard to find sweet point)

Hough transform - experiments




Courtesy of TTK Automation Technology Laboratory

Generalized Hough transform


D. Ballard, Generalizing the Hough Transform to Detect Arbitrary Shapes, Pattern Recognition 13(2), 1981

- Identify a shape model by measuring the location of its parts and shape centroid
- Measurements: orientation θ , location of p
- Each measurement casts a vote in the Hough space: $p + r(\theta)$




Generalized Hough transform

B. Leibe, A. Leonardis, and B. Schiele, [Combined Object Categorization and Segmentation with an Implicit Shape Model](#), ECCV Workshop on Statistical Learning in Computer Vision 2004



Hough transform



Courtesy of unknown

Figure 5.2 (a) An image containing two lines, sampled irregularly, and several random points. (b) Plot of the counters in the corresponding parameter space (how many points contribute to each cell (m, n)). Notice that the main peaks are obvious, but there are many secondary peaks.

Same cons and pros as before...

QUESTIONS?