

---

# Computer Vision & Image Processing

## CSE 473 / 573

Instructor - Kevin R. Keane, PhD

TAs - Radhakrishna Dasari, Yuhao Du, Niyazi Sorkunlu

**Lecture 26**  
**October 30, 2017**  
**Image Alignment**

# Image Alignment

---

Readings for today:

- Forsyth and Ponce chapter 12

# Image alignment

---







Image from [http://graphics.cs.cmu.edu/courses/15-463/2010\\_fall/](http://graphics.cs.cmu.edu/courses/15-463/2010_fall/)

# Alignment applications

---

- [A look into the past](#)



# Alignment applications

---

- [A look into the past](#)





# Alignment applications

---

- Streetside images



# Alignment applications

---

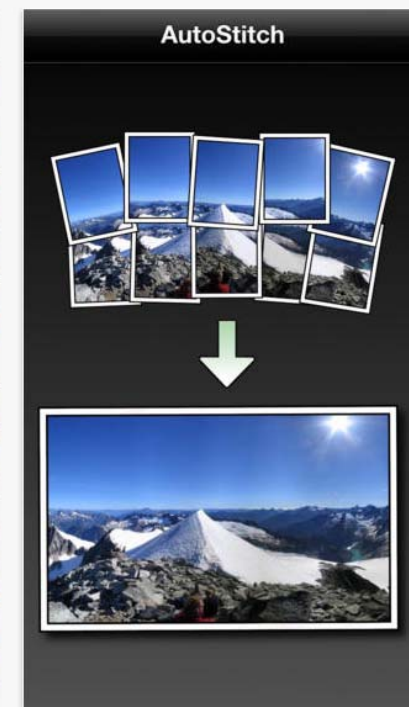
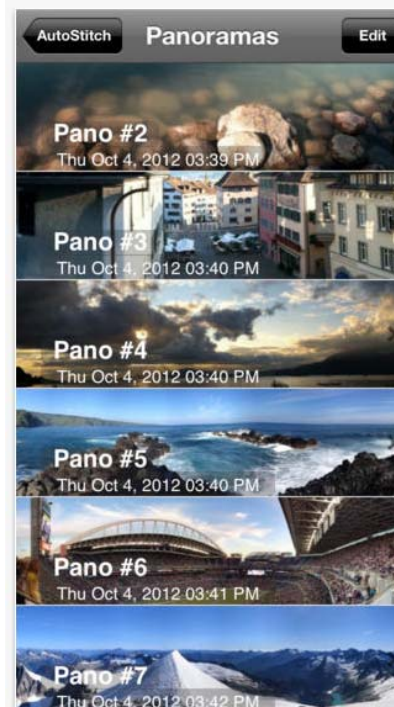


Panorama stitching

**AutoStitch Panorama**  
By Cloudburst Research Inc.  
Open iTunes to buy and download apps.



[View In iTunes](#)



# Alignment applications

---




Recognition  
of object  
instances



# Alignment challenges

---



Small degree of overlap  
Intensity changes




Occlusion,  
clutter

# Feature-based alignment

---

- Search sets of feature matches that agree in terms of:
  - a) Local appearance
  - b) Geometric configuration



# Feature-based alignment: Overview


---

- Alignment as fitting
  - Affine transformations
  - Homographies
- Robust alignment
  - Descriptor-based feature matching
  - RANSAC
- Large-scale alignment
  - Inverted indexing
  - Vocabulary trees
- Application: searching the night sky

# Alignment as fitting

---

- Previous lectures: fitting a model to features in one image



Find model  $M$  that minimizes


$$\sum_i \text{residual}(x_i, M)$$



# Alignment as fitting

---



- Previous lectures: fitting a model to features in one image



Find model  $M$  that minimizes

$$\sum_i \text{residual}(x_i, M)$$

- Alignment: fitting a model to a transformation between pairs of features (*matches*) in two images




Find transformation  $T$  that minimizes

$$\sum_i \text{residual}(T(x_i), x'_i)$$


# 2D transformation models

---


- Similarity  
(translation, scale, rotation)



- Affine




- Projective  
(homography)



# Let's start with affine transformations

---


- Simple fitting procedure (linear least squares)
- Approximates viewpoint changes for roughly planar objects and roughly orthographic cameras
- Can be used to initialize fitting for more complex models



# Fitting an affine transformation

---

- Assume we know the correspondences, how do we get the transformation?



$$\mathbf{x}'_i = \mathbf{M}\mathbf{x}_i + \mathbf{t}$$

$$\begin{bmatrix} x'_i \\ y'_i \end{bmatrix} = \begin{bmatrix} m_1 & m_2 \\ m_3 & m_4 \end{bmatrix} \begin{bmatrix} x_i \\ y_i \end{bmatrix} + \begin{bmatrix} t_1 \\ t_2 \end{bmatrix}$$

Want to find  $\mathbf{M}$ ,  $\mathbf{t}$  to minimize


$$\sum_{i=1}^n \| \mathbf{x}'_i - \mathbf{M}\mathbf{x}_i - \mathbf{t} \|^2$$



# Fitting an affine transformation

---

- Assume we know the correspondences, how do we get the transformation?




$$\begin{bmatrix} x'_i \\ y'_i \end{bmatrix} = \begin{bmatrix} m_1 & m_2 \\ m_3 & m_4 \end{bmatrix} \begin{bmatrix} x_i \\ y_i \end{bmatrix} + \begin{bmatrix} t_1 \\ t_2 \end{bmatrix}$$
$$\begin{bmatrix} m_1 \\ m_2 \\ m_3 \\ m_4 \\ t_1 \\ t_2 \end{bmatrix} = \begin{bmatrix} \dots \\ x'_i \\ y'_i \\ \dots \end{bmatrix}$$

# Fitting an affine transformation

---

$$\begin{bmatrix} \dots & & & & & \\ x_i & y_i & 0 & 0 & 1 & 0 \\ 0 & 0 & x_i & y_i & 0 & 1 \\ \dots & & & & & \end{bmatrix} \begin{bmatrix} m_1 \\ m_2 \\ m_3 \\ m_4 \\ t_1 \\ t_2 \end{bmatrix} = \begin{bmatrix} \dots \\ x'_i \\ y'_i \\ \dots \end{bmatrix}$$


- Linear system with six unknowns
- Each match gives us two linearly independent equations: need at least three to solve for the transformation parameters



# Fitting a plane projective transformation

---

- **Homography:** plane projective transformation  
(transformation taking a quad to another arbitrary quad)



# Homography

---

- The transformation between two views of a planar surface



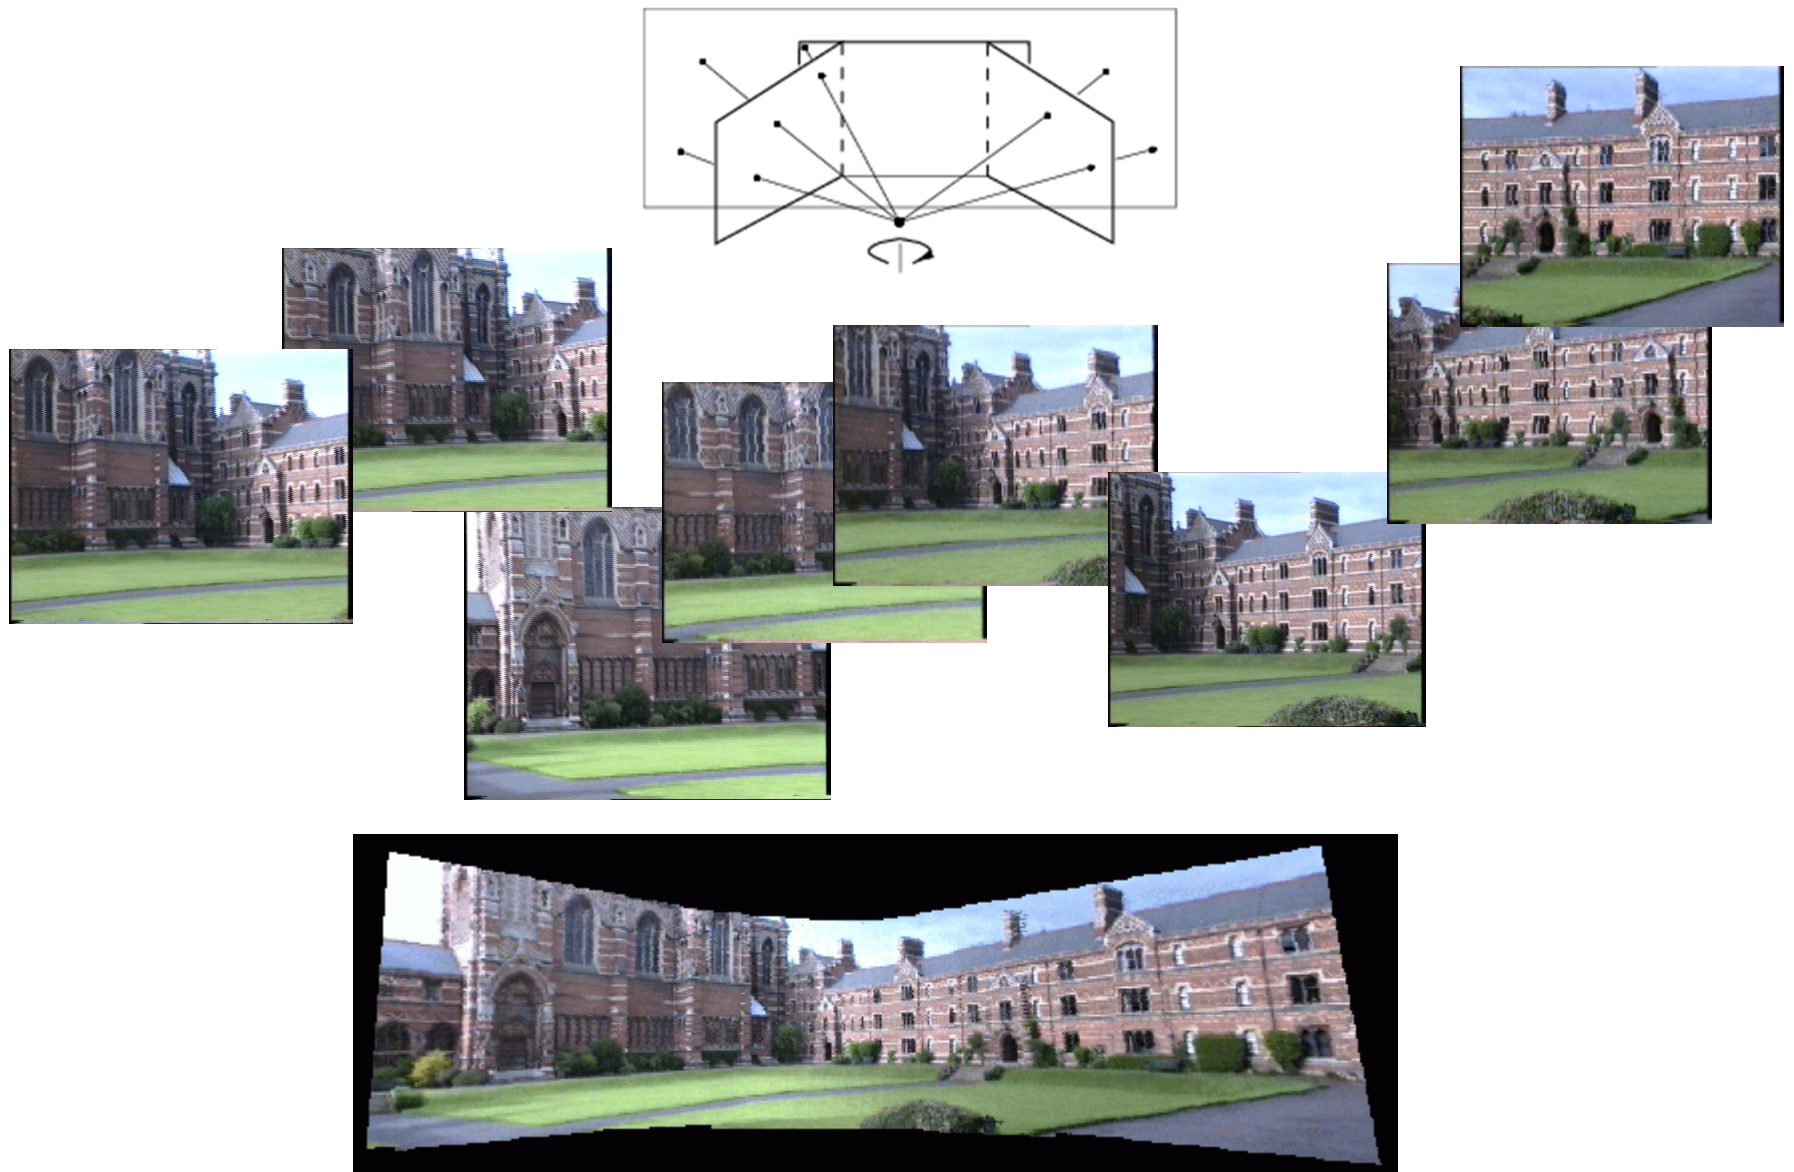
- The transformation between images from two cameras that share the same center





# Application: Panorama stitching

---



Source: Hartley & Zisserman

# Fitting a homography

---

- Recall: homogeneous coordinates

$$(x, y) \Rightarrow \begin{bmatrix} x \\ y \\ 1 \end{bmatrix}$$

Converting *to* homogeneous  
image coordinates

$$\begin{bmatrix} x \\ y \\ w \end{bmatrix} \Rightarrow (x/w, y/w)$$

Converting *from* homogeneous  
image coordinates

# Fitting a homography

---

- Recall: homogeneous coordinates

$$(x, y) \Rightarrow \begin{bmatrix} x \\ y \\ 1 \end{bmatrix}$$

Converting *to* homogeneous  
image coordinates

$$\begin{bmatrix} x \\ y \\ w \end{bmatrix} \Rightarrow (x/w, y/w)$$

Converting *from* homogeneous  
image coordinates

- Equation for homography:

$$\lambda \begin{bmatrix} x' \\ y' \\ 1 \end{bmatrix} = \begin{bmatrix} h_{11} & h_{12} & h_{13} \\ h_{21} & h_{22} & h_{23} \\ h_{31} & h_{32} & h_{33} \end{bmatrix} \begin{bmatrix} x \\ y \\ 1 \end{bmatrix}$$

# Fitting a homography

---

- Equation for homography:

$$\lambda \begin{bmatrix} x'_i \\ y'_i \\ 1 \end{bmatrix} = \begin{bmatrix} h_{11} & h_{12} & h_{13} \\ h_{21} & h_{22} & h_{23} \\ h_{31} & h_{32} & h_{33} \end{bmatrix} \begin{bmatrix} x_i \\ y_i \\ 1 \end{bmatrix} \quad \begin{aligned} \lambda \mathbf{x}'_i &= \mathbf{H} \mathbf{x}_i \\ \mathbf{x}'_i \times \mathbf{H} \mathbf{x}_i &= 0 \end{aligned}$$

$$\begin{bmatrix} x'_i \\ y'_i \\ 1 \end{bmatrix} \times \begin{bmatrix} \mathbf{h}_1^T \mathbf{x}_i \\ \mathbf{h}_2^T \mathbf{x}_i \\ \mathbf{h}_3^T \mathbf{x}_i \end{bmatrix} = \begin{bmatrix} y'_i \mathbf{h}_3^T \mathbf{x}_i - \mathbf{h}_2^T \mathbf{x}_i \\ \mathbf{h}_1^T \mathbf{x}_i - x'_i \mathbf{h}_3^T \mathbf{x}_i \\ x'_i \mathbf{h}_2^T \mathbf{x}_i - y'_i \mathbf{h}_1^T \mathbf{x}_i \end{bmatrix}$$

$$\begin{bmatrix} 0^T & -\mathbf{x}_i^T & y'_i \mathbf{x}_i^T \\ \mathbf{x}_i^T & 0^T & -x'_i \mathbf{x}_i^T \\ -y'_i \mathbf{x}_i^T & x'_i \mathbf{x}_i^T & 0^T \end{bmatrix} \begin{pmatrix} \mathbf{h}_1 \\ \mathbf{h}_2 \\ \mathbf{h}_3 \end{pmatrix} = 0 \quad \begin{aligned} &3 \text{ equations,} \\ &\text{only 2 linearly} \\ &\text{independent} \end{aligned}$$



# Fitting a homography

---


$$\begin{bmatrix} 0^T & \mathbf{x}_1^T & -y'_1 \mathbf{x}_1^T \\ \mathbf{x}_1^T & 0^T & -x'_1 \mathbf{x}_1^T \\ \dots & \dots & \dots \\ 0^T & \mathbf{x}_n^T & -y'_n \mathbf{x}_n^T \\ \mathbf{x}_n^T & 0^T & -x'_n \mathbf{x}_n^T \end{bmatrix} \begin{pmatrix} \mathbf{h}_1 \\ \mathbf{h}_2 \\ \mathbf{h}_3 \end{pmatrix} = 0 \quad \mathbf{A} \mathbf{h} = 0$$

- H has 8 degrees of freedom (9 parameters, but scale is arbitrary)
- One match gives us two linearly independent equations
- Homogeneous least squares: find  $\mathbf{h}$  minimizing  $\|\mathbf{A}\mathbf{h}\|^2$ 
  - Eigenvector of  $\mathbf{A}^T\mathbf{A}$  corresponding to smallest eigenvalue
  - Four matches needed for a minimal solution

# Robust feature-based alignment

---


- So far, we've assumed that we are given a set of “ground-truth” correspondences between the two images we want to align
- What if we don't know the correspondences?



# Robust feature-based alignment

---

- So far, we've assumed that we are given a set of "ground-truth" correspondences between the two images we want to align
- What if we don't know the correspondences?





# Robust feature-based alignment

---



# Robust feature-based alignment

---



- Extract features



# Robust feature-based alignment


---



- Extract features
- Compute *putative matches*

# Robust feature-based alignment


---



- Extract features
- Compute *putative matches*
- Loop:
  - *Hypothesize* transformation  $T$

# Robust feature-based alignment

---



- Extract features
- Compute *putative matches*
- Loop:
  - *Hypothesize* transformation  $T$
  - *Verify* transformation (search for other matches consistent with  $T$ )



# Robust feature-based alignment


---



- Extract features
- Compute *putative matches*
- Loop:
  - *Hypothesize* transformation  $T$
  - *Verify* transformation (search for other matches consistent with  $T$ )

# Generating putative correspondences

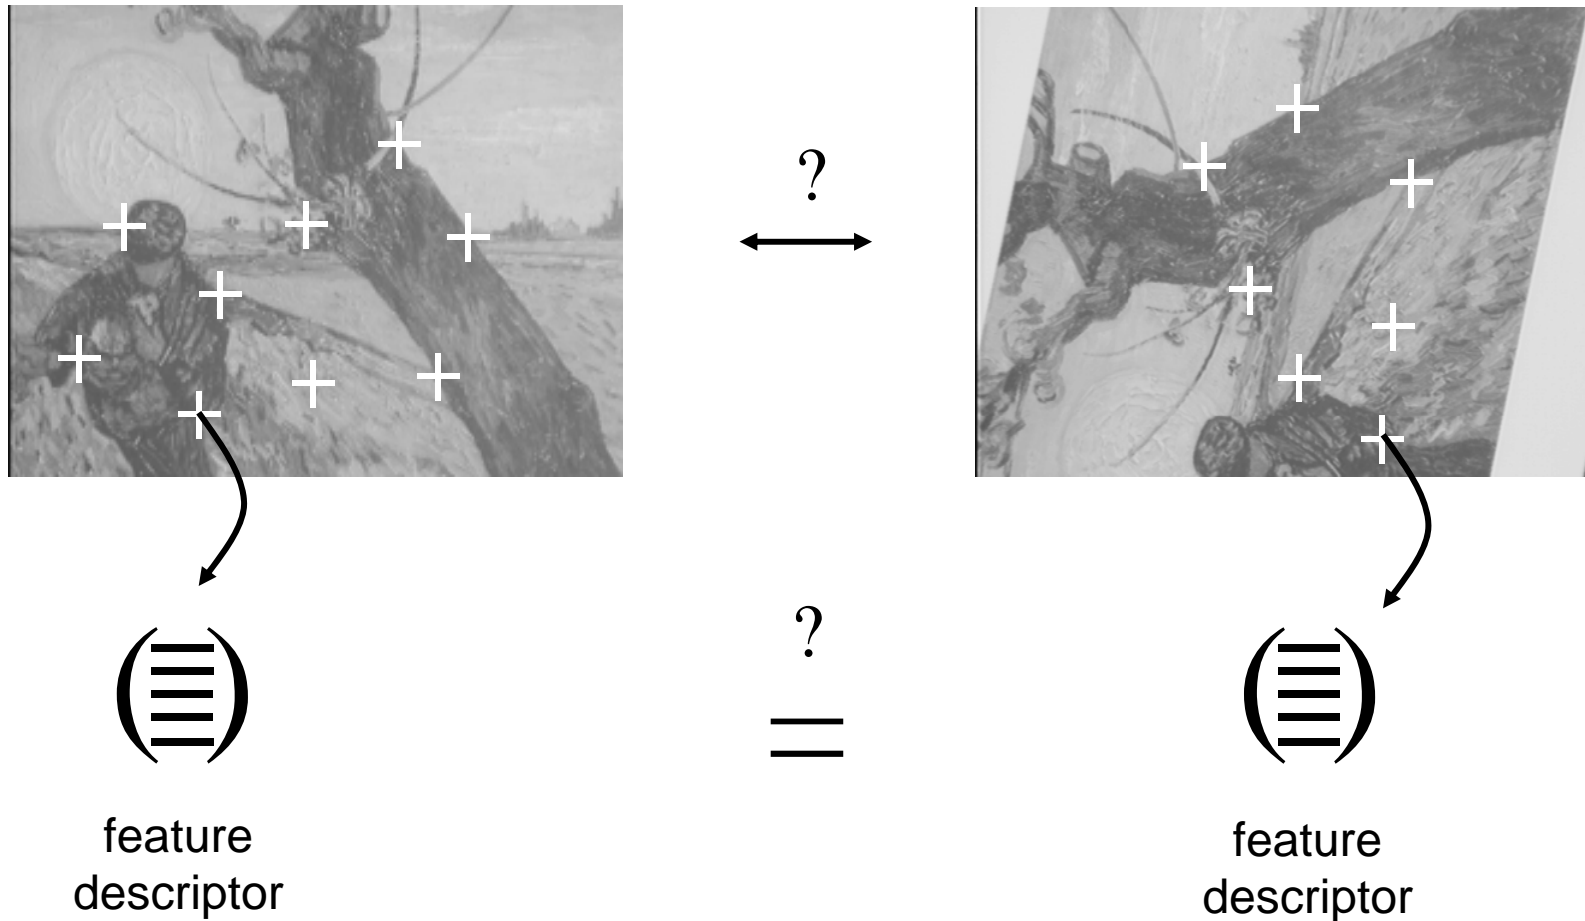
---





# Generating putative correspondences

---

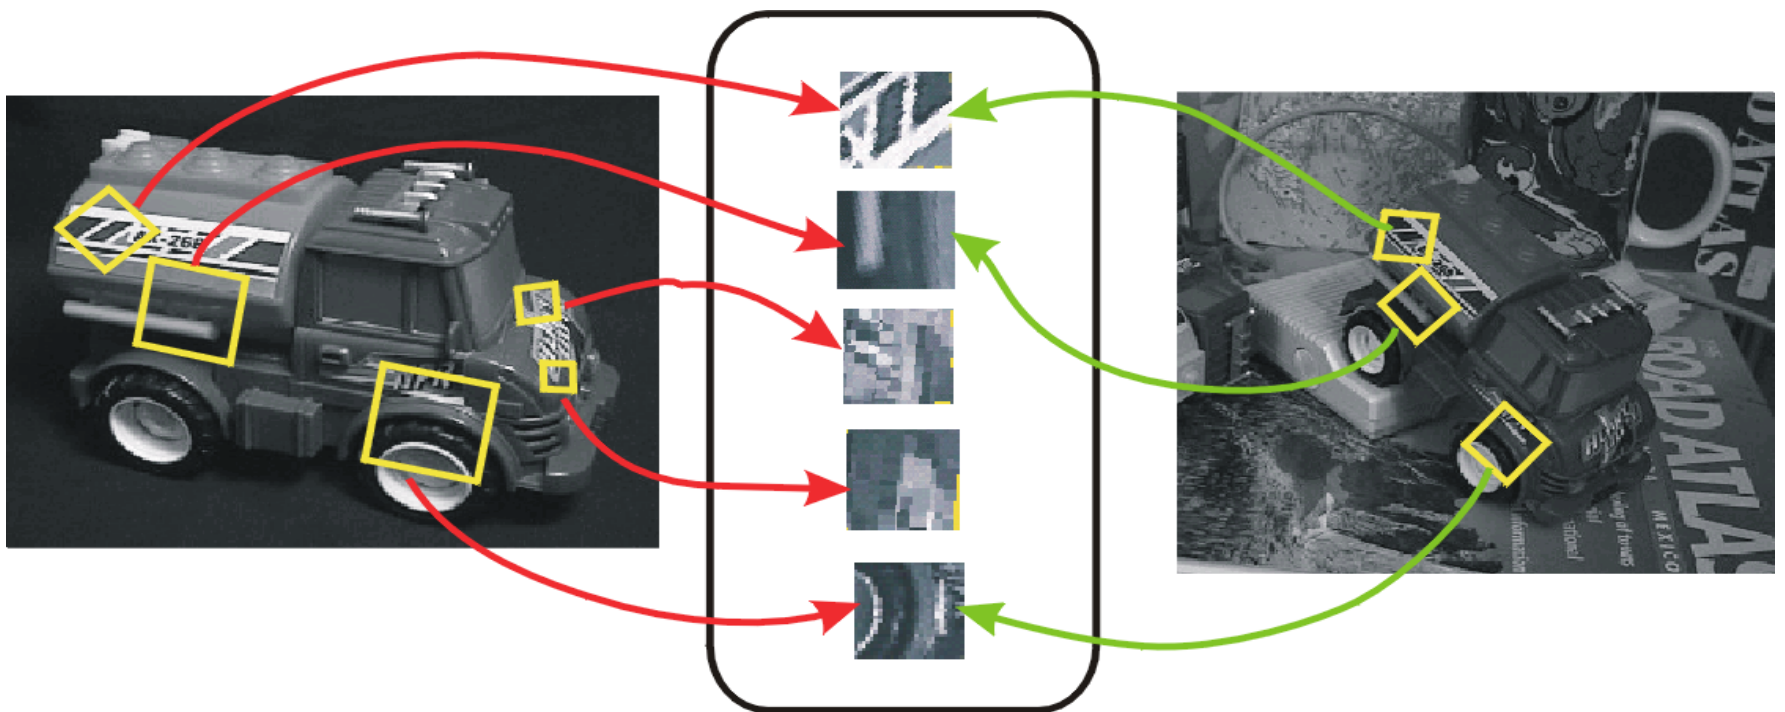


- Need to compare *feature descriptors* of local patches surrounding interest points

# Feature descriptors

---

- Recall: feature detection and description



# Feature descriptors

---

- Simplest descriptor: vector of raw intensity values
- How to compare two such vectors?
  - Sum of squared differences (SSD)

$$\text{SSD}(\mathbf{u}, \mathbf{v}) = \sum_i (u_i - v_i)^2$$

– Not invariant to intensity change


- Normalized correlation

$$\rho(\mathbf{u}, \mathbf{v}) = \frac{(\mathbf{u} - \bar{\mathbf{u}}) \cdot (\mathbf{v} - \bar{\mathbf{v}})}{\|\mathbf{u} - \bar{\mathbf{u}}\| \|\mathbf{v} - \bar{\mathbf{v}}\|} = \frac{\sum_i (u_i - \bar{\mathbf{u}})(v_i - \bar{\mathbf{v}})}{\sqrt{\left(\sum_j (u_j - \bar{\mathbf{u}})^2\right) \left(\sum_j (v_j - \bar{\mathbf{v}})^2\right)}}$$

– Invariant to affine intensity change

## Disadvantage of intensity vectors as descriptors


- Small deformations can affect the matching score a lot



# Feature descriptors: SIFT

---

- Descriptor computation:
  - Divide patch into 4x4 sub-patches
  - Compute histogram of gradient orientations (8 reference angles) inside each sub-patch
  - Resulting descriptor:  $4 \times 4 \times 8 = 128$  dimensions



David G. Lowe. ["Distinctive image features from scale-invariant keypoints."](#) *IJCV* 60 (2), pp. 91-110, 2004.

# Feature descriptors: SIFT

---

- Descriptor computation:
  - Divide patch into 4x4 sub-patches
  - Compute histogram of gradient orientations (8 reference angles) inside each sub-patch
  - Resulting descriptor:  $4 \times 4 \times 8 = 128$  dimensions
- Advantage over raw vectors of pixel values
  - Gradients less sensitive to illumination change
  - Pooling of gradients over the sub-patches achieves robustness to small shifts, but still preserves some spatial information


David G. Lowe. ["Distinctive image features from scale-invariant keypoints."](#) *IJCV* 60 (2), pp. 91-110, 2004.





# Feature matching

---

- Generating *putative matches*: for each patch in one image, find a short list of patches in the other image that could match it based solely on appearance




# Problem: Ambiguous putative matches



# Rejection of unreliable matches

---

- How can we tell which putative matches are more reliable?
- Heuristic: compare distance of **nearest** neighbor to that of **second** nearest neighbor
  - Ratio of closest distance to second-closest distance will be *high* for features that are *not* distinctive



**Threshold of 0.8  
provides good  
separation**

David G. Lowe. ["Distinctive image features from scale-invariant keypoints."](#) *IJCV* 60 (2), pp. 91-110, 2004.

# RANSAC

---

- The set of putative matches contains a very high percentage of outliers


## **RANSAC loop:**

1. Randomly select a *seed group* of matches
2. Compute transformation from seed group
3. Find *inliers* to this transformation
4. If the number of inliers is sufficiently large, re-compute least-squares estimate of transformation on all of the inliers

Keep the transformation with the largest number of inliers

# RANSAC example: Translation


---





# RANSAC example: Translation


---





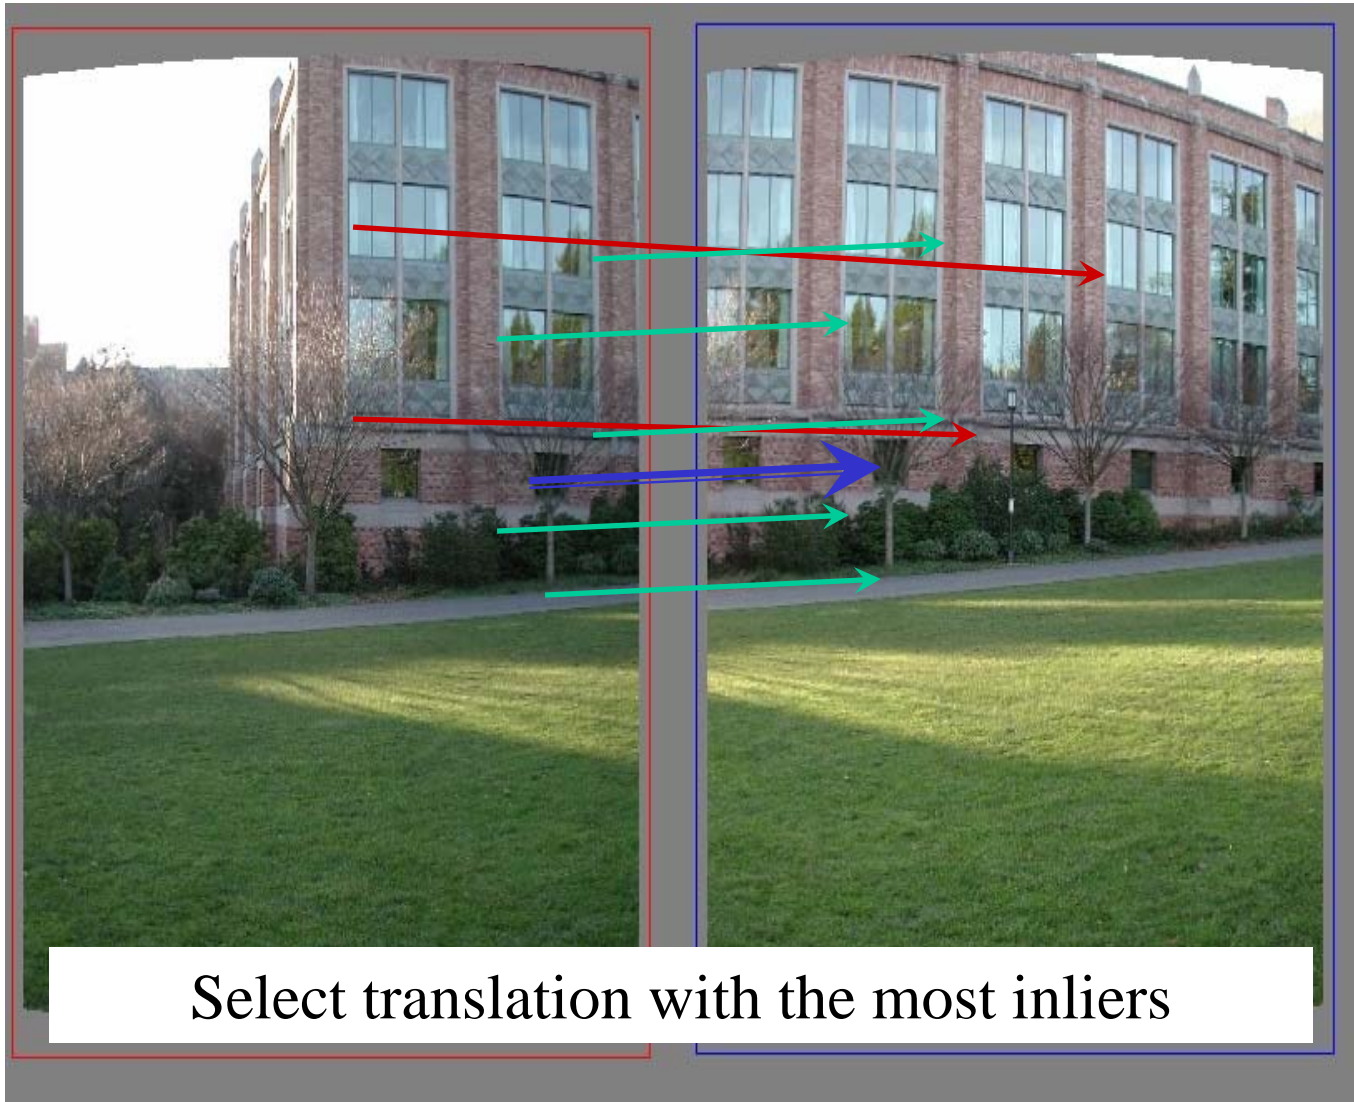
# RANSAC example: Translation

---



# RANSAC example: Translation



---



# Scalability: Alignment to large databases

---

- What if we need to align a test image with thousands or millions of images in a model database?
  - Efficient putative match generation
    - Approximate descriptor similarity search, inverted indices



Test image



# Large-scale visual search

---

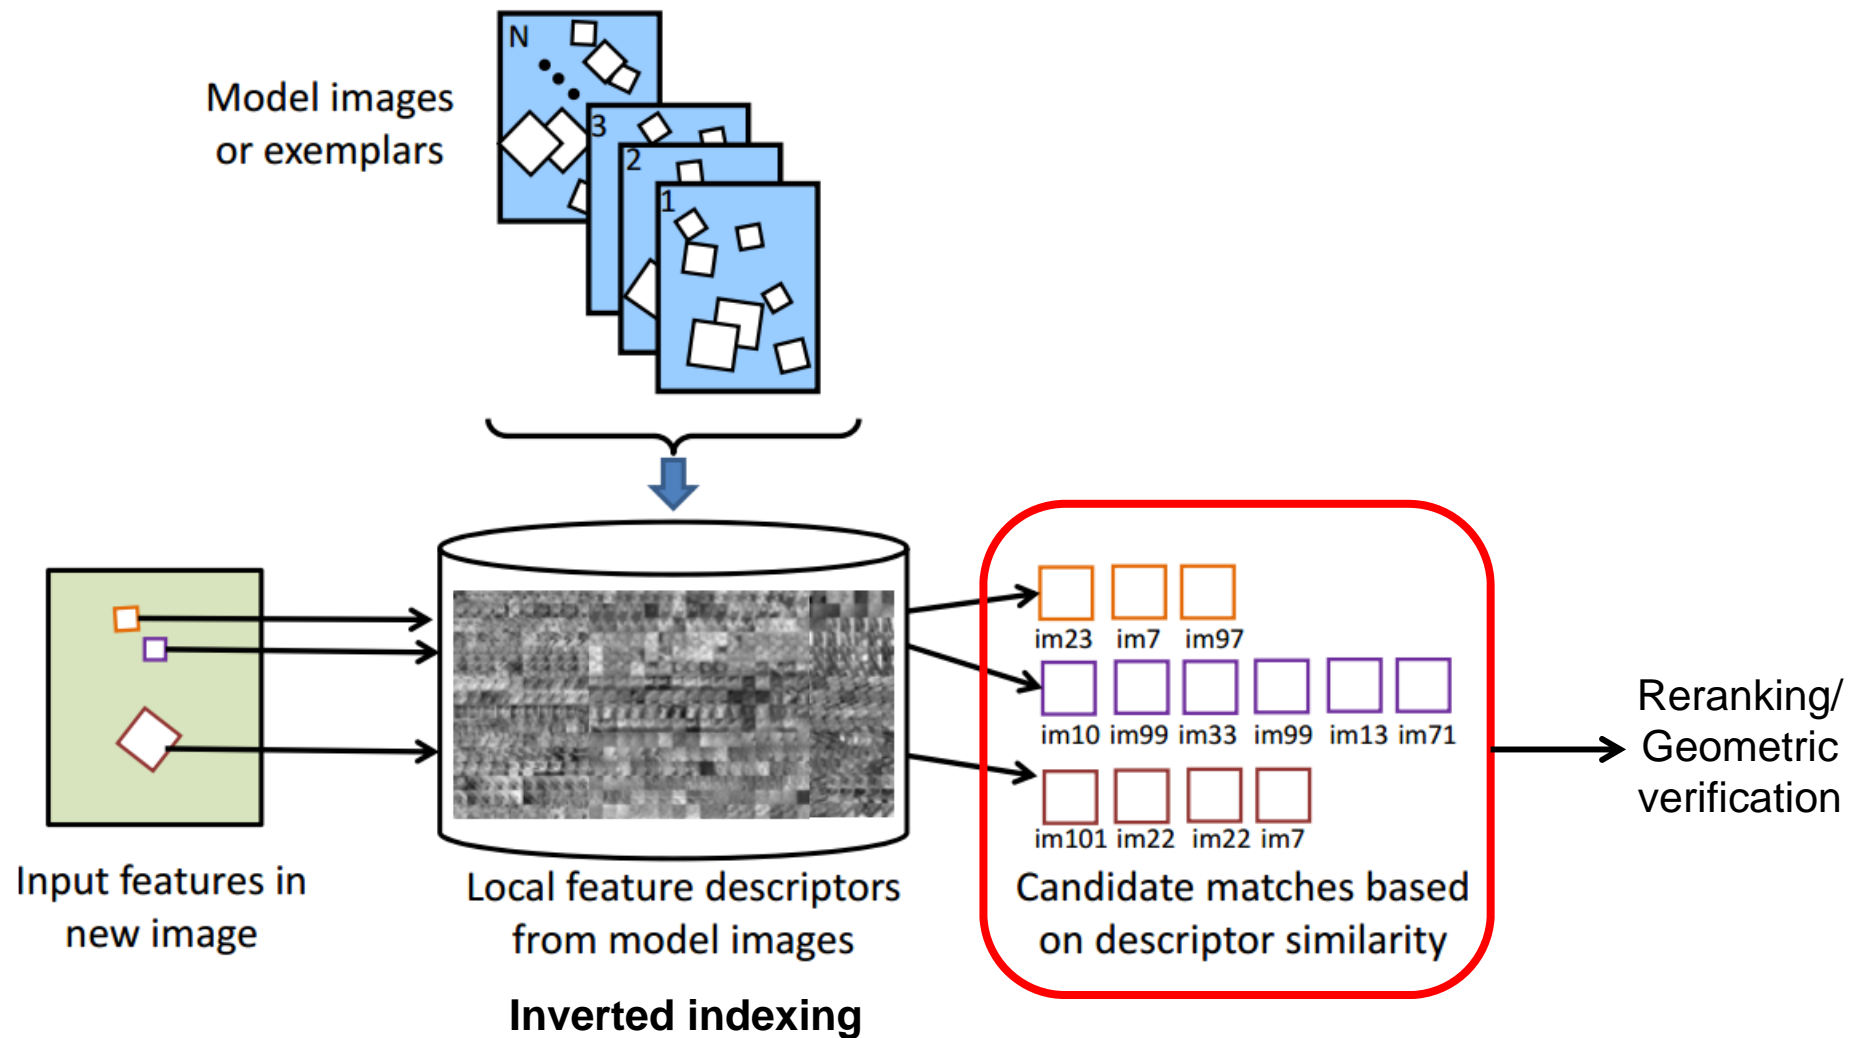
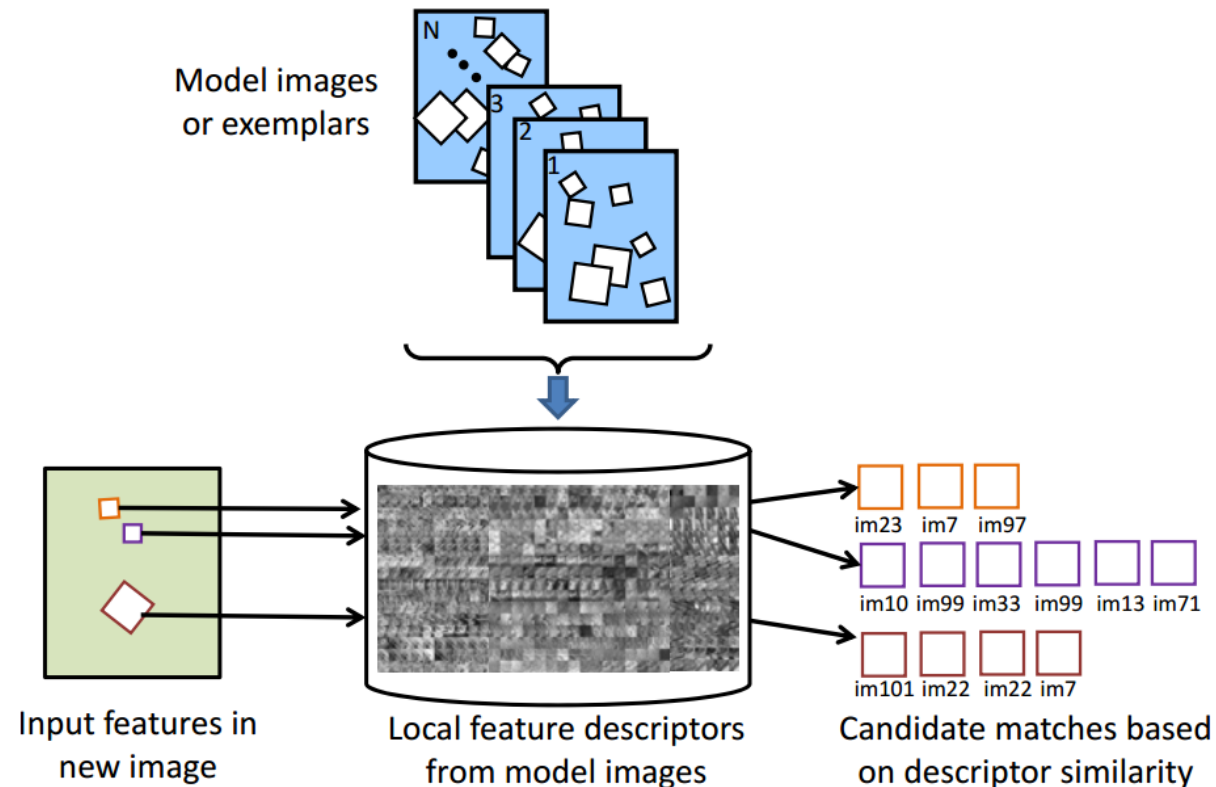


Figure from: Kristen Grauman and Bastian Leibe, [Visual Object Recognition](#), Synthesis Lectures on Artificial Intelligence and Machine Learning, April 2011, Vol. 5, No. 2, Pages 1-181



# How to do the indexing?


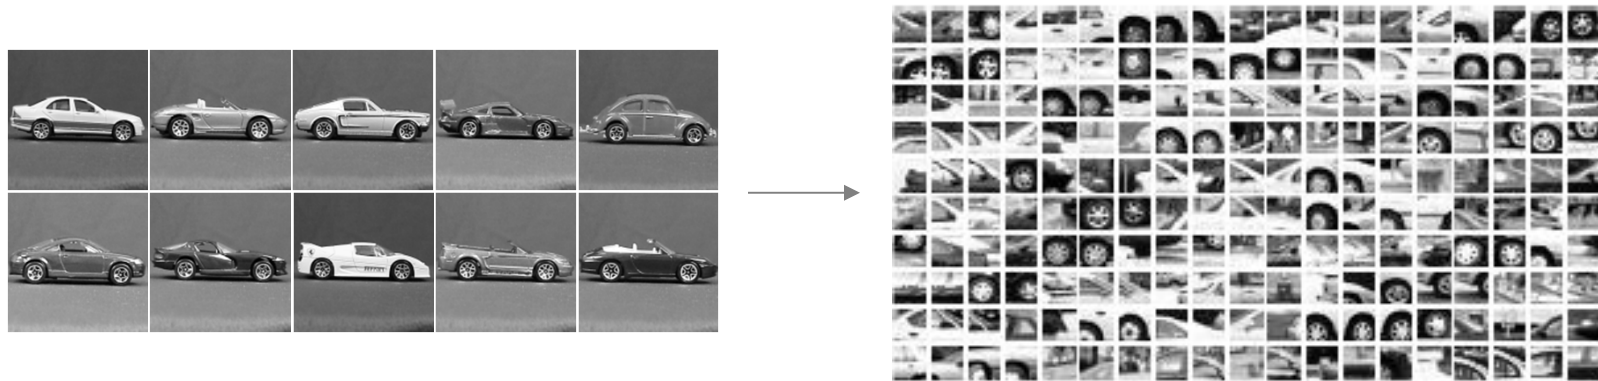
---



- Idea: find a set of *visual codewords* to which descriptors can be *quantized*

# Recall: Visual codebook for generalized Hough transform

---





# K-means clustering

---

- Want to minimize sum of squared Euclidean distances between points  $\mathbf{x}_i$  and their nearest cluster centers  $\mathbf{m}_k$

$$D(X, M) = \sum_{\text{cluster } k} \sum_{\substack{\text{point } i \text{ in} \\ \text{cluster } k}} (\mathbf{x}_i - \mathbf{m}_k)^2$$

Algorithm:

- Randomly initialize K cluster centers
- Iterate until convergence:
  - Assign each data point to the nearest center
  - Recompute each cluster center as the mean of all points assigned to it

# K-means demo

---

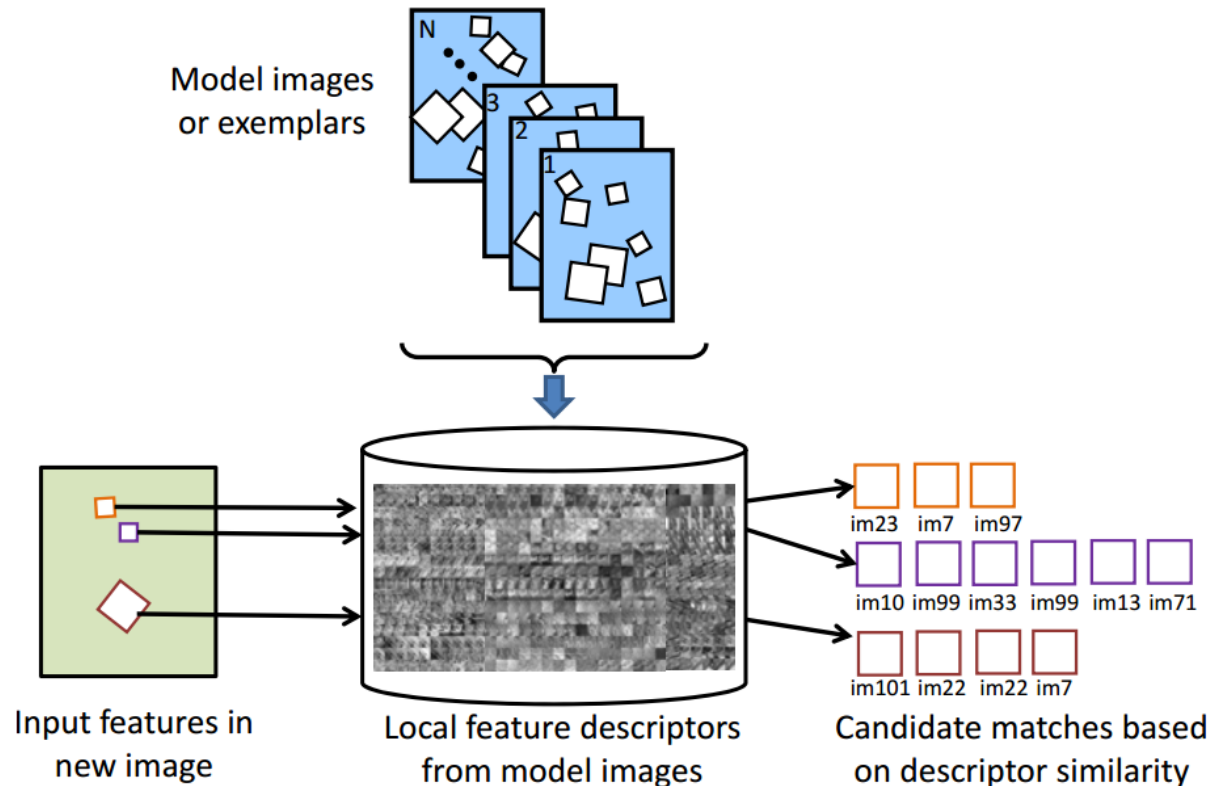


Source: <http://shabal.in/visuals/kmeans/1.html>

Another demo: <http://www.kovan.ceng.metu.edu.tr/~maya/kmeans/>

# How to do the indexing?

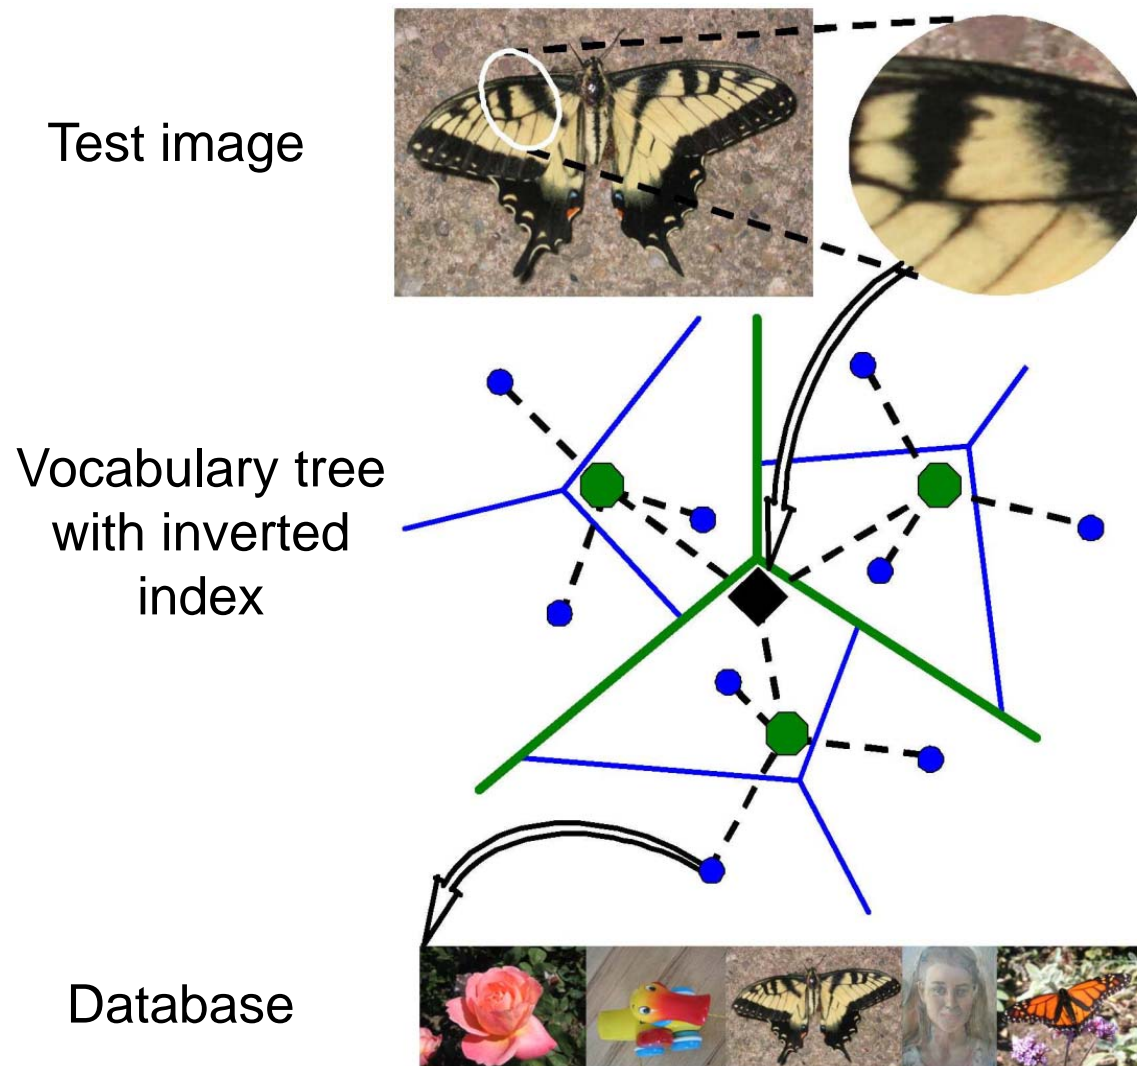
---




- Cluster descriptors in the database to form codebook
- At query time, quantize descriptors in query image to nearest codevectors
- Problem solved?

# Efficient indexing technique: Vocabulary trees


---





Hierarchical k-means  
clustering of  
descriptor space  
(vocabulary tree)

Slide credit: D. Nister




Vocabulary tree/inverted index

Slide credit: D. Nister




Model images



Populating the vocabulary tree/inverted index

Slide credit: D. Nister

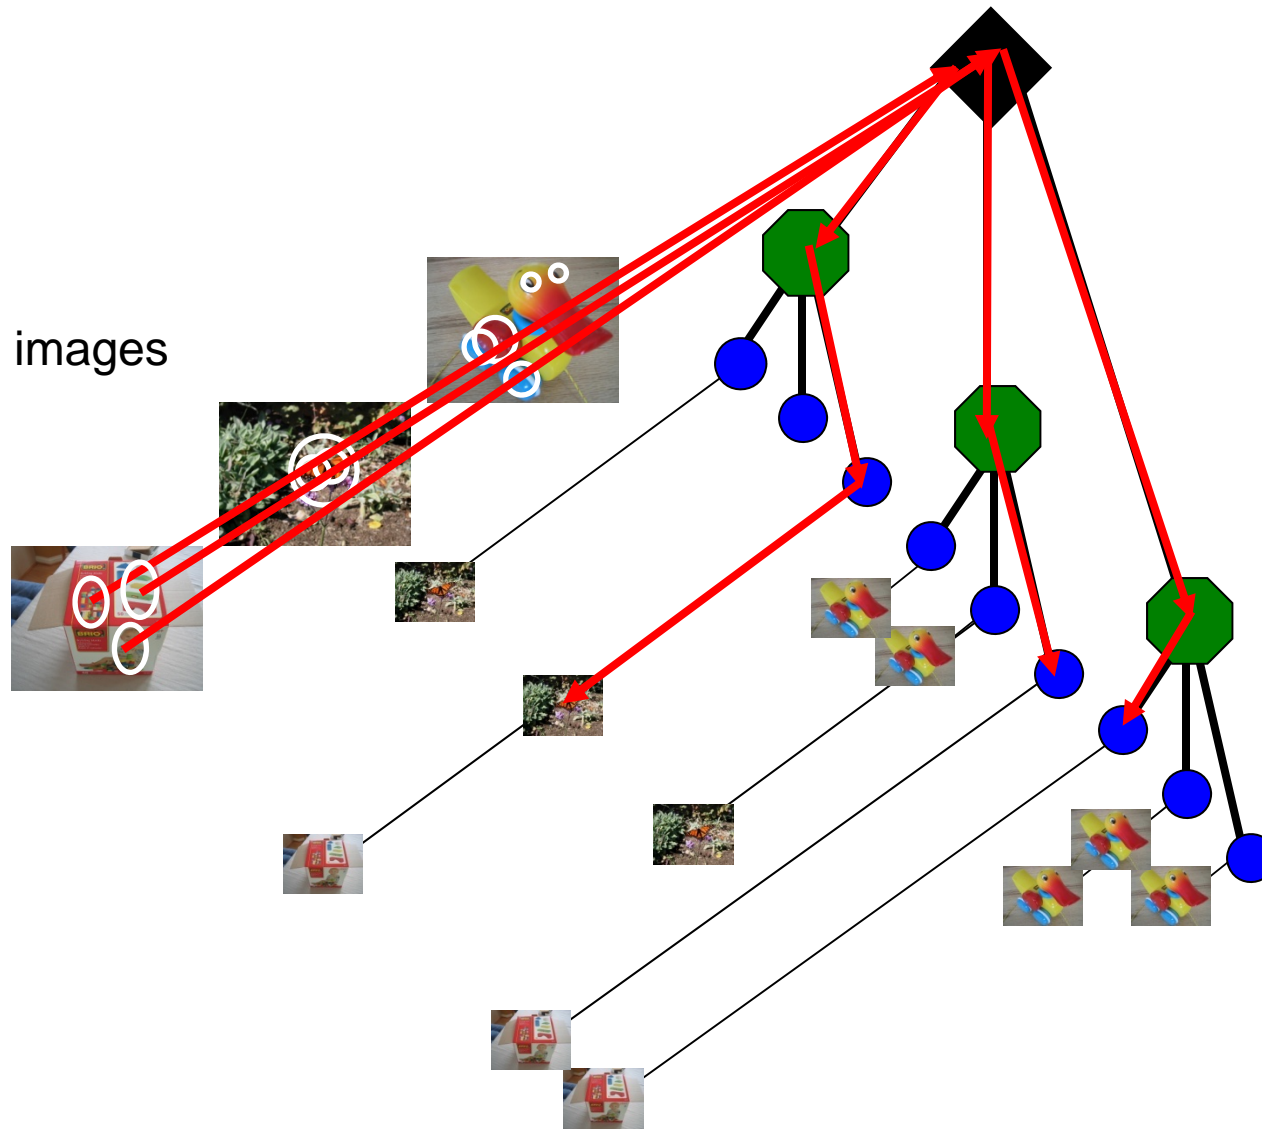
Model images



Populating the vocabulary tree/inverted index

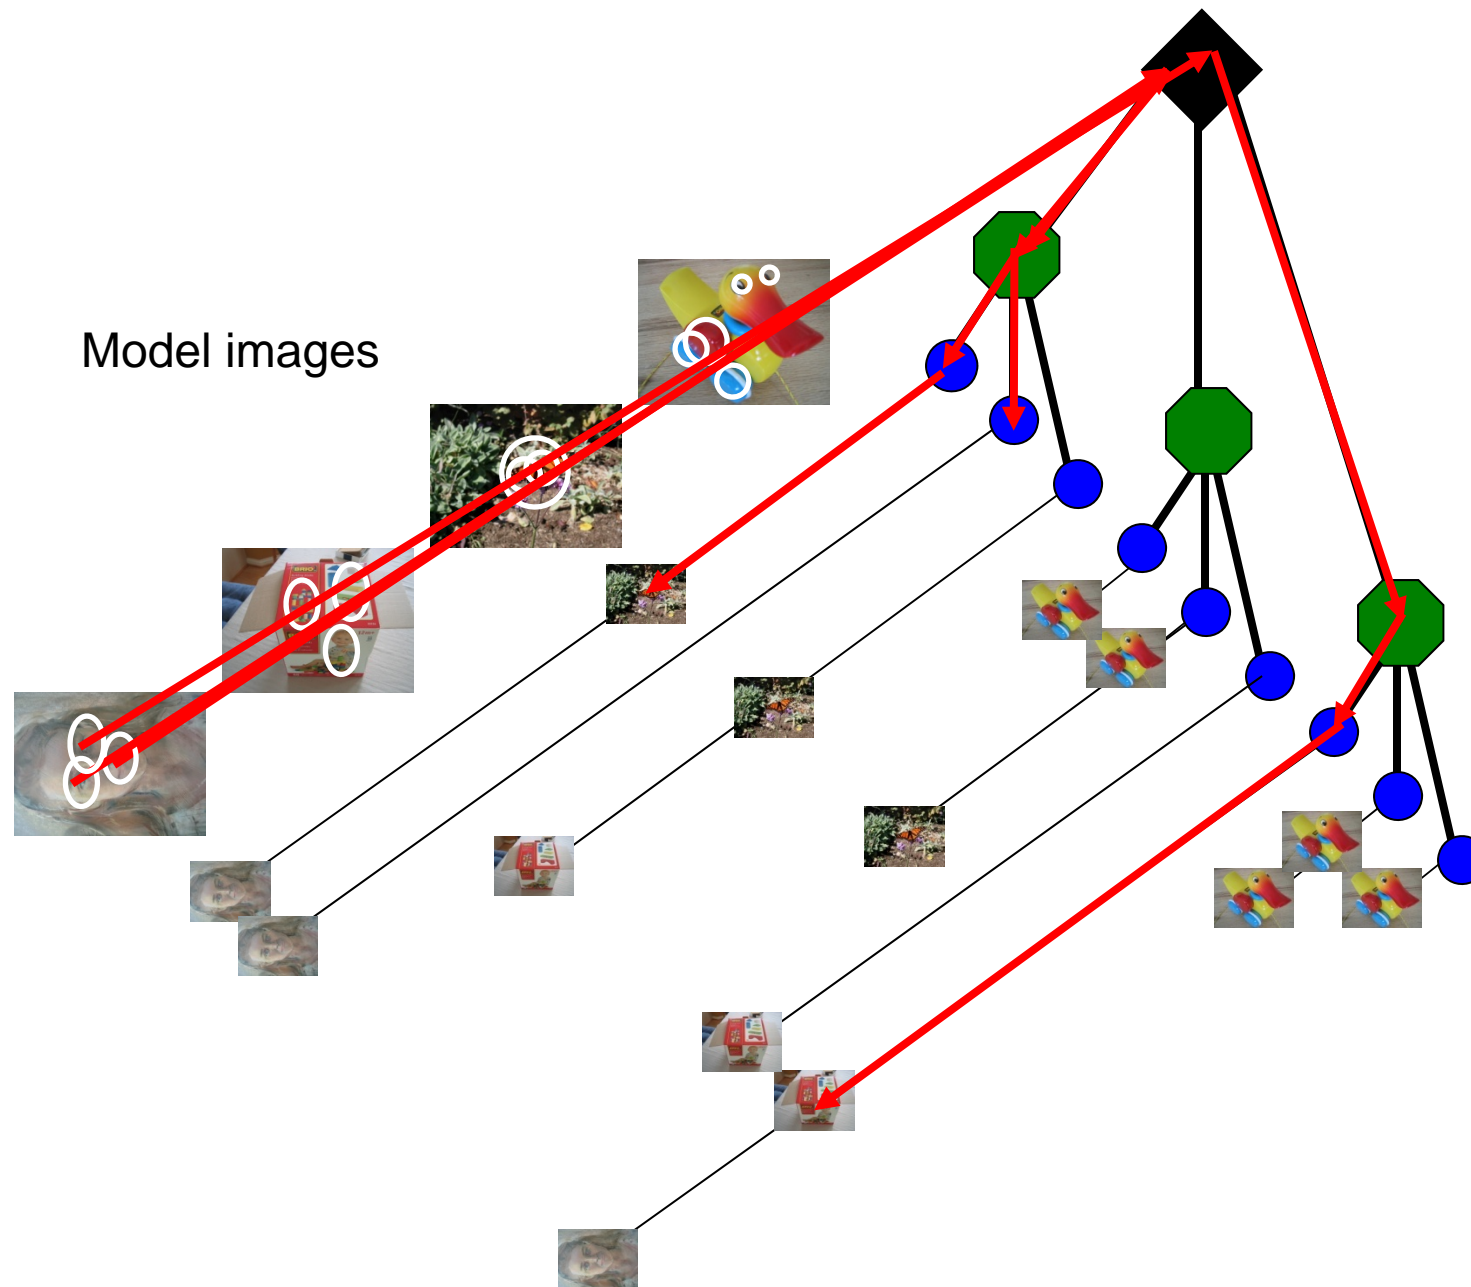
Slide credit: D. Nister

Model images



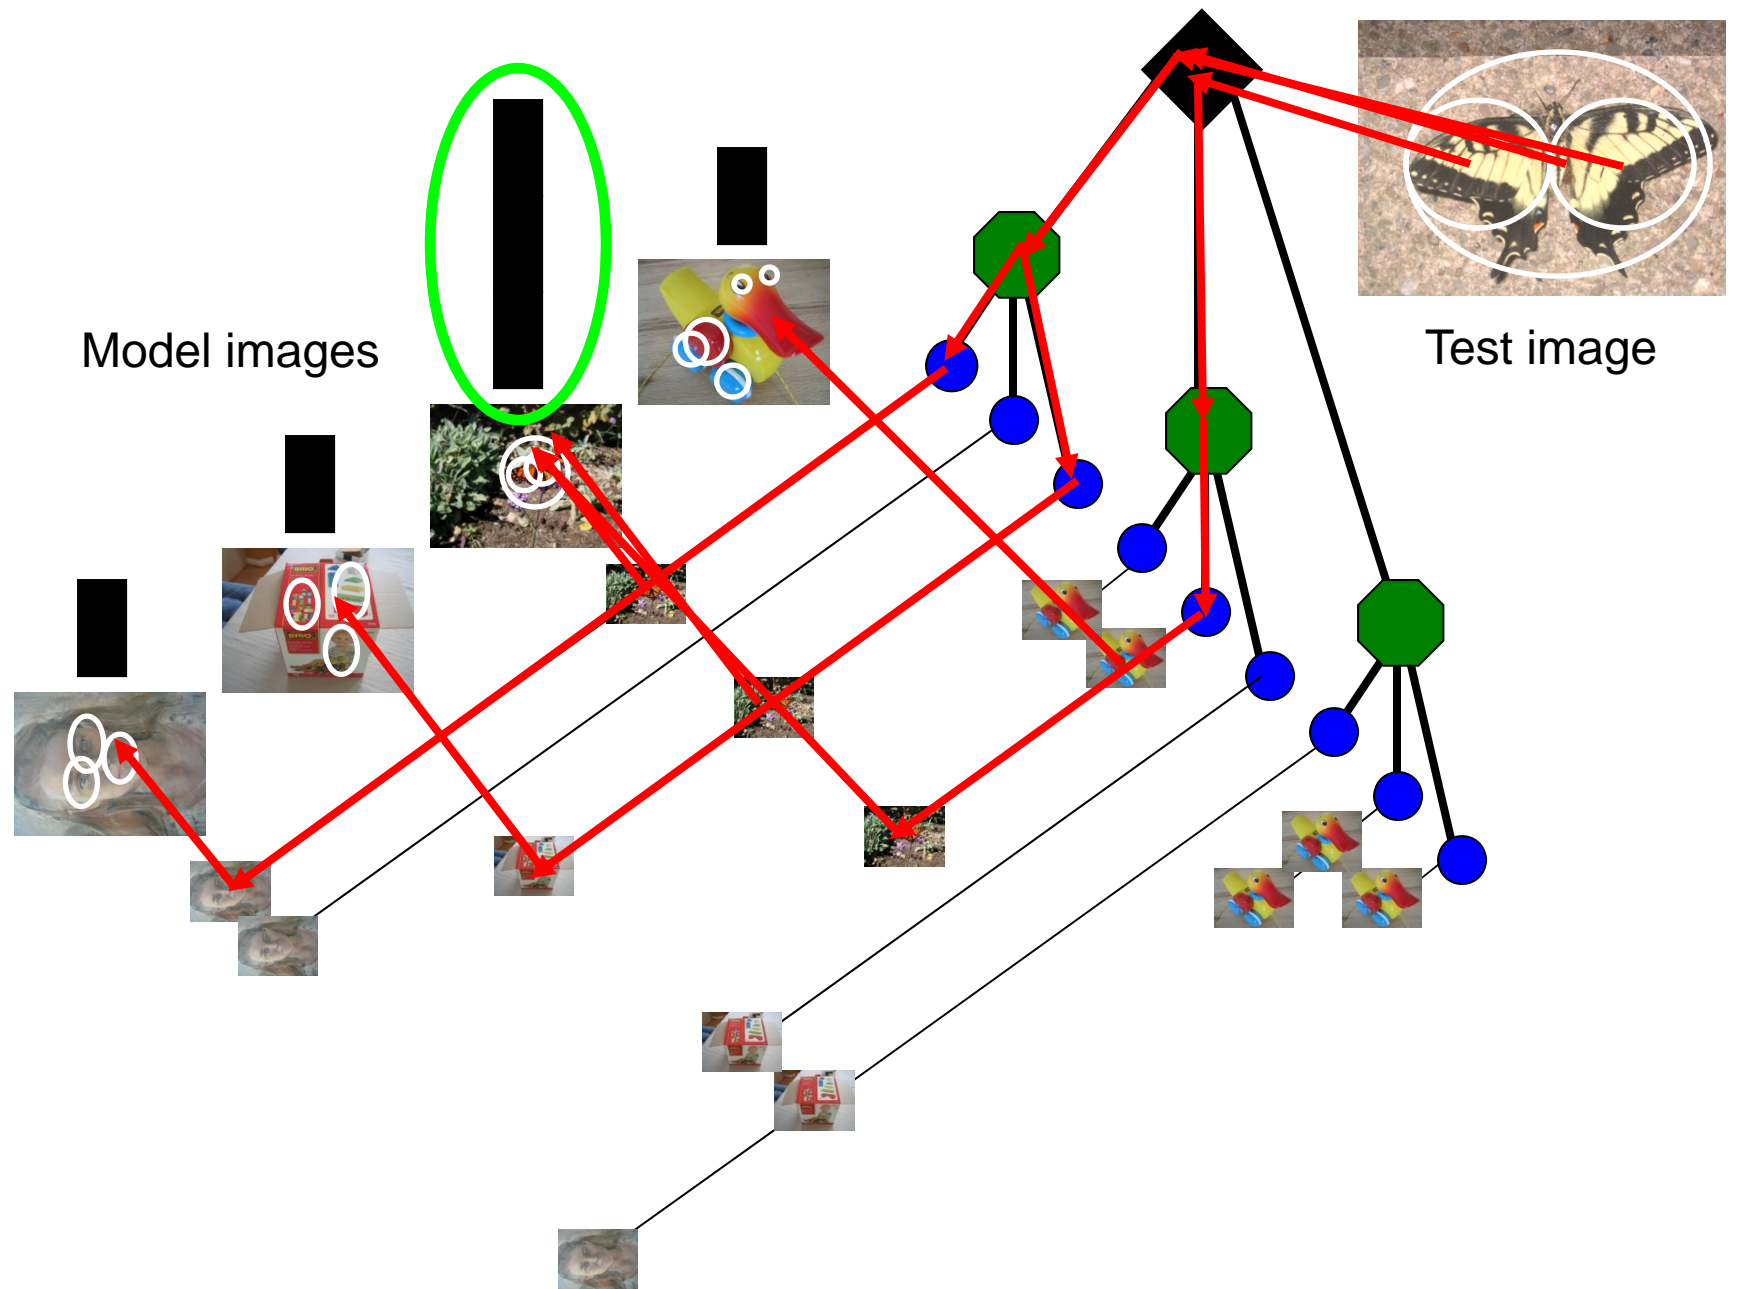
Populating the vocabulary tree/inverted index

Slide credit: D. Nister



Populating the vocabulary tree/inverted index

Slide credit: D. Nister




Looking up a test image

Slide credit: D. Nister

Cool application of large-scale alignment:  
searching the night sky

<http://www.astrometry.net/>





# Slide Credits

---

Svetlana Lazebnik – UIUC

Derek Hoiem – UIUC

David Forsyth - UIUC

---

# Questions



# Image Alignment

---

Readings for today:

- Forsyth and Ponce chapter 12